

LAMP: Linear Verification of Matrix Multiplication via Proximity Testing

Anonymous Submission

Abstract—Verifiable computation systems often need to prove large matrix multiplication statements, but a direct SNARK arithmetization of a $k \times k$ product requires $\mathcal{O}(k^3)$ constraints. Freivalds’ randomized check reduces the algebraic computation to vector-matrix products, but proving those products inside a SNARK still costs $\mathcal{O}(k^2)$ constraints.

We present **LAMP**, a matrix-multiplication checking protocol that combines Freivalds’ randomized check with proximity testing over linear error-correcting codes. The prover commits to encoded matrices and intermediate vectors before the sampled query positions are derived. The SNARK circuit then checks only the sampled codeword positions, while Merkle openings and CP-Link proofs bind the vector witnesses and sampled circuit values to the committed values. We prove soundness for this committed-input setting under the soundness of the SNARK backend, the binding of the commitments, the correctness of the CP-Link checks, and the distance of the code.

For a fixed number t of sampled positions, the main in-circuit SNARK relation has $\mathcal{O}(tk)$ constraints, with additional $\mathcal{O}(t \log n) + E_{\text{link}}(k)$ backend work for Merkle openings and CP-Link checks. We implement **LAMP** in Go and compare it with a Freivalds-based SNARK circuit. In the matrix benchmark at $k = 2^{12}$, **LAMP** reduces the constraint count by $30.3\times$ and shortens proof generation time by $7.24\times$; verification stays below 0.01 seconds in the measured range.

Index Terms—Verifiable computation, SNARKs, Proximity testing, Linear codes, Freivalds’ algorithm, Merkle trees, Commitment linking

1. Introduction

Large-scale matrix multiplication is a recurring bottleneck in verifiable computation. A prover may wish to convince a verifier that matrices $\mathbf{A}, \mathbf{B}, \mathbf{C}$ satisfy $\mathbf{AB} = \mathbf{C}$, without requiring either the verifier or the SNARK circuit to recompute the entire product. This problem appears throughout verifiable inference and training for modern neural networks. In particular, projection and attention layers often involve matrix multiplications with dimension $k \geq 1,000$. A naïve arithmetization of a single $k \times k$ matrix multiplication inside a SNARK circuit generates $\mathcal{O}(k^3)$ arithmetic constraints, which is impractical even for moderate values of k .

The $\mathcal{O}(k^2)$ barrier. Freivalds’ classical randomized check [1] reduces this cost from $\mathcal{O}(k^3)$ to $\mathcal{O}(k^2)$. Given matrices $\mathbf{A}, \mathbf{B}, \mathbf{C}$ and a random vector $\mathbf{r}$, one verifies

$$\mathbf{rAB} = \mathbf{rC}$$

instead of recomputing the full product. Several recent systems, including Mystique [2], zkML [3], and the Freivalds-based baseline in [4], adopt this strategy by directly encoding the resulting vector-matrix products as SNARK constraints. The resulting $\mathcal{O}(k^2)$ -size circuit is an improvement, but it remains costly. In our optimized baseline, Freivalds already requires 50.5 million R1CS constraints and 405.04s of proving time at $k = 4,096$. Thus, a significant gap remains between the $\mathcal{O}(k^2)$ SNARK cost and practical large-scale matrix verification.

Our insight: From computation to checking. We observe that the $\mathcal{O}(k^2)$ barrier is not inherent to the verification problem itself. Rather, it arises because existing approaches *compute* vector-matrix products inside the SNARK circuit. Our key idea is to shift the paradigm from “computing inside the circuit” to “*checking* inside the circuit.” The prover performs the expensive matrix arithmetic outside the SNARK and commits to the relevant intermediate values. The verifier then checks only a small number of lightweight algebraic relations that are sufficient to detect inconsistencies.

The technical core of this shift is *proximity testing over linear error-correcting codes*. By encoding each matrix row-wise using a systematic linear code $\mathcal{C}$ with relative distance δ , the vector-matrix products

$$\mathbf{x} = \mathbf{rA}, \quad \mathbf{y} = \mathbf{xB}, \quad \mathbf{z} = \mathbf{rC}$$

can be checked column by column in the encoded domain. The linearity of $\mathcal{C}$ ensures that a single inner product, or “fold,” over one codeword column is sufficient to check one coordinate of the encoded product. If the prover cheats in any of these relations, the corresponding codewords disagree on at least a δ -fraction of positions, and a random query detects the inconsistency with probability at least δ .

The key point is that the SNARK circuit does not take the full matrices as witnesses. Instead, it receives only t sampled columns and checks local algebraic relations on those columns, including fold computations and the codeword encoding consistency of the intermediate vectors $\mathbf{x}, \mathbf{y}$, and $\mathbf{z}$. We use a commit-and-prove SNARK to produce commitments to the full vector witnesses and sampled values used inside the circuit. Together with the SNARK proof, the prover sends Merkle membership proofs

for the queried indices in the previously committed Merkle trees, as well as CP-Link proofs showing that the pre-query vector commitments, the corresponding Merkle leaf values, and the SNARK witnesses represent the same underlying data. By verifying these proofs, the verifier is assured that the values used inside the circuit are consistent with the data committed before the SNARK proof was generated. As a result, the in-circuit checking cost is $\mathcal{O}(tk)$, while the additional opening and linking work is $\mathcal{O}(t \log n) + E_{\text{link}}(k)$. Since t is the number of sampled codeword positions and is independent of the matrix dimension k , the protocol has linear in-circuit cost in the matrix dimension.

Contributions. We present LAMP, a protocol for efficiently proving matrix multiplication statements. Our contributions are as follows.

- **A code-based protocol for matrix multiplication.** We propose LAMP, which combines Freivalds’ randomized reduction with proximity testing over linear error-correcting codes. Instead of computing vector-matrix products inside the SNARK circuit, LAMP encodes the matrices and intermediate vectors as codewords and checks only randomly sampled codeword positions. By exploiting the linearity of the code, each sampled position can be verified through local fold relations, reducing the in-circuit checking cost from $\mathcal{O}(k^2)$ to $\mathcal{O}(tk)$, which is linear in the matrix dimension for fixed t .
- **Soundness analysis.** We prove the soundness of LAMP. In particular, we show that if $\mathbf{AB} \neq \mathbf{C}$, then any prover that commits to inconsistent encoded values is detected with probability determined by the code distance and the number of sampled positions. The analysis also accounts for the Merkle openings and CP-links that bind the values checked inside the SNARK circuit to the data committed before proof generation.
- **Implementation and evaluation.** We implement LAMP in Go and evaluate it against a Freivalds-based SNARK circuit for matrix dimensions from $k = 2^7$ to 2^{13} . At $k = 2^{12}$, LAMP reduces the constraint count by $30.3\times$ and shortens proof generation time by $7.24\times$, while verification remains independent of the matrix dimension and takes less than 0.01 s. We also evaluate LAMP on a GPT-2 workload with sequence length 2^{10} , where it reduces the constraint count by $2.3\times$ and proving time by $1.78\times$, showing that the approach extends to matrix-heavy neural network workloads.

2. Preliminaries

We introduce notation and recall the building blocks used in our construction. Throughout, $\mathbb{F}$ denotes a finite field of size $|\mathbb{F}| = p$ for a prime p , and λ denotes the security parameter. We write $\text{negl}(\lambda)$ for any function that is negligible in λ , and PPT for probabilistic polynomial-time.

2.1. Notation

Matrices are denoted by bold uppercase letters $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{k \times k}$, and vectors by bold lowercase letters $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{F}^k$. For a vector $\mathbf{x}$, we write $\mathbf{x}[i]$ for its i -th entry. For a matrix $\mathbf{M}$, we write $\mathbf{M}[i, j]$ for the entry in the i -th row and j -th column. We denote the i -th row of $\mathbf{M}$ by $\mathbf{M}[i, *]$, and we will often write it simply as $\mathbf{M}[i]$. Similarly, we denote the j -th column of $\mathbf{M}$ by $\mathbf{M}[:, j]$. We write $[n] = \{1, 2, \dots, n\}$. For $t \in \mathbb{N}$, $I \leftarrow \$ [n]^t$ denotes the uniform sampling of t indices from $[n]$. Similarly, $r \leftarrow \$ \mathbb{F}$ denotes uniform random sampling from $\mathbb{F}$.

2.2. Linear Error-Correcting Codes

A linear code over $\mathbb{F}$ is an $\mathbb{F}$ -linear map $\mathcal{C} : \mathbb{F}^k \rightarrow \mathbb{F}^n$ defined by a generator matrix $\mathbf{G} \in \mathbb{F}^{k \times n}$, so that $\mathcal{C}(\mathbf{m}) = \mathbf{m} \cdot \mathbf{G}$ for any input vector $\mathbf{m} \in \mathbb{F}^k$. The code has rate $\rho = k/n$ and minimum distance $d = \min_{\mathbf{c} \neq \mathbf{c}'} \Delta(\mathbf{c}, \mathbf{c}')$, where the minimum is taken over distinct codewords $\mathbf{c}, \mathbf{c}' \in \mathcal{C}$, and Δ denotes the Hamming distance. The relative distance is $\delta = d/n$. A code is systematic if the first k coordinates of $\mathcal{C}(\mathbf{m})$ equal $\mathbf{m}$ itself.

We write $\text{Enc}(\mathbf{m}) := \mathcal{C}(\mathbf{m})$ for the encoding of an input vector. For a matrix $\mathbf{M} \in \mathbb{F}^{k \times k}$, we define the row-wise encoding $\hat{\mathbf{M}} := \text{Enc}(\mathbf{M}) \in \mathbb{F}^{k \times n}$ as the matrix whose i -th row is $\text{Enc}(\mathbf{M}[i, *])$ for each $i \in [k]$. The j -th column of $\hat{\mathbf{M}}$ is denoted by $\hat{\mathbf{M}}[:, j] \in \mathbb{F}^k$ for each $j \in [n]$.

The key property we exploit is *linearity*: for any coefficients $r_1, \dots, r_k \in \mathbb{F}$,

$$\sum_{i=1}^k r_i \cdot \hat{\mathbf{M}}[i, *] = \text{Enc}\left(\sum_{i=1}^k r_i \cdot \mathbf{M}[i, *]\right). \quad (1)$$

Definition 2.1 (Fold). For vectors $\mathbf{c}, \mathbf{r} \in \mathbb{F}^k$ with $\mathbf{c} = (c_0, \dots, c_{k-1})^\top$ and $\mathbf{r} = (r_0, \dots, r_{k-1})$, define

$$\text{Fold}(\mathbf{c}, \mathbf{r}) := \langle \mathbf{r}, \mathbf{c} \rangle = \sum_{i=0}^{k-1} r_i \cdot c_i.$$

Equivalently, for the j -th column of a row-wise encoded matrix, $\text{Fold}(\hat{\mathbf{M}}[:, j], \mathbf{r})$ equals the j -th coordinate of $\text{Enc}(\mathbf{rM})$. This follows directly from the linearity of the code (Equation 1).

Concrete instantiations. Our protocol is compatible with any linear code. Two natural choices are: (i) *Reed–Solomon codes*, where $\mathcal{C}(\mathbf{m})$ evaluates the polynomial $\sum_i m_i X^i$ over a domain of size $n > k$, achieving the optimal (MDS) distance $\delta = 1 - k/n + 1/n$; and (ii) *expander-based codes* as used in Brakedown [5], which achieve linear-time encoding $\mathcal{O}(n)$ with constant relative distance. We discuss the trade-offs in Section 4.3.

2.3. Commitment Schemes

A commitment scheme allows a committer to bind itself to a message while keeping the message hidden until it is

opened. In this work, we use two types of commitments: Merkle tree commitments and Pedersen vector commitments.

2.3.1. Merkle tree commitments. A Merkle tree commitment over a collision-resistant hash function $H : \{0, 1\}^* \rightarrow \{0, 1\}^\lambda$ is a vector commitment for a vector $\mathbf{v} = (\mathbf{v}[1], \dots, \mathbf{v}[n]) \in \mathbb{F}^n$. It consists of the following algorithms.

- $\text{rt} \leftarrow \text{MT.Com}(\mathbf{v})$: outputs a root hash rt obtained by constructing a binary hash tree with leaves $H(\mathbf{v}[1]), \dots, H(\mathbf{v}[n])$.
- $(\mathbf{v}[j], \pi_j) \leftarrow \text{MT.Open}(\mathbf{v}, j)$: outputs the leaf value $\mathbf{v}[j]$ and its authentication path π_j for position $j \in [n]$.
- $b \leftarrow \text{MT.Verify}(\text{rt}, j, \mathbf{v}[j], \pi_j)$: outputs $b = 1$ if $\mathbf{v}[j]$ is accepted as the j -th leaf consistent with root rt , and outputs $b = 0$ otherwise.

Merkle tree commitments are *position-binding*: under the collision resistance of H , no PPT adversary can produce a root rt , an index $j \in [n]$, two distinct values $v \neq v'$, and authentication paths π_j, π'_j such that

$$\text{MT.Verify}(\text{rt}, j, v, \pi_j) = \text{MT.Verify}(\text{rt}, j, v', \pi'_j) = 1,$$

except with probability $\epsilon_{\text{merkle}}(\lambda) = \text{negl}(\lambda)$.

2.3.2. Pedersen commitments. Pedersen commitments can be viewed as algebraic vector commitments over a group $\mathbb{G}$ of prime order q . $\text{ck} = (G_1, \dots, G_n, H)$ be a commitment key consisting of fixed generators in $\mathbb{G}$. A Pedersen commitment is defined by the following algorithms.

- $\text{cm} \leftarrow \text{Ped.Com}(\text{ck}, \mathbf{v}; r)$: outputs a commitment $\text{cm} := \sum_{i=1}^n \mathbf{v}[i]G_i + rH \in \mathbb{G}$ for a vector $\mathbf{v} \in \mathbb{F}^n$ and randomness $r \leftarrow \mathbb{F}$.
- $\pi \leftarrow \text{Ped.Open}(\mathbf{v}, r)$: outputs the opening $\pi = (\mathbf{v}, r)$ consisting of the committed vector and the randomness.
- $b \leftarrow \text{Ped.Verify}(\text{ck}, \text{cm}, \mathbf{v}, r)$: outputs $b = 1$ if $\text{cm} \stackrel{?}{=} \sum_{i=1}^n \mathbf{v}[i]G_i + rH$, and outputs $b = 0$ otherwise.

Pedersen commitments are *perfectly hiding* and *computationally binding*: under the discrete logarithm assumption, no PPT adversary can produce a commitment cm , two distinct vectors $\mathbf{v} \neq \mathbf{v}'$, and randomness values r, r' such that

$$\text{Ped.Verify}(\text{ck}, \text{cm}, \mathbf{v}, r) = \text{Ped.Verify}(\text{ck}, \text{cm}, \mathbf{v}', r') = 1,$$

except with probability $\epsilon_{\text{ped}}(\lambda) = \text{negl}(\lambda)$.

They are also *additively homomorphic*: for $\text{cm}_1 = \text{Ped.Com}(\text{ck}, \mathbf{v}_1; r_1)$ and $\text{cm}_2 = \text{Ped.Com}(\text{ck}, \mathbf{v}_2; r_2)$, $\text{cm}_1 + \text{cm}_2 = \text{Ped.Com}(\text{ck}, \mathbf{v}_1 + \mathbf{v}_2; r_1 + r_2)$.

2.4. Succinct Non-interactive Arguments of Knowledge

A SNARK (succinct non-interactive argument of knowledge) for an NP relation $\mathcal{R}$ is a tuple of algorithms $\Pi = (\text{Setup}, \text{Prove}, \text{Verify})$.

- $\text{pp} \leftarrow \Pi.\text{Setup}(1^\lambda, \mathcal{R})$: generates public parameters pp for the relation $\mathcal{R}$.
- $\pi \leftarrow \Pi.\text{Prove}(\text{pp}, \mathbf{x}, \mathbf{w})$: outputs a proof π for a statement $\mathbf{x}$ and witness $\mathbf{w}$ satisfying $(\mathbf{x}, \mathbf{w}) \in \mathcal{R}$.
- $b \leftarrow \Pi.\text{Verify}(\text{pp}, \mathbf{x}, \pi)$: outputs $b = 1$ if π is accepted as a valid proof for $\mathbf{x}$, and outputs $b = 0$ otherwise.

2.5. CP-Link

A *CP-Link* proof is a lightweight commit-and-prove proof whose only purpose is to link two commitments to the same value. The value may be a field element, a vector, or a matrix column. For a SNARK-side commitment Γ and an external commitment C , the link relation is

$$\mathcal{R}_{\text{link}} = \{ ((\Gamma, C), (u, \rho_\Gamma, \rho_C)) : \Gamma = \text{Com}_{\text{snark}}(u; \rho_\Gamma) \wedge C = \text{Com}_{\text{ext}}(u; \rho_C) \}.$$

We use this relation in two places. First, for each sampled position, it links a SNARK-side commitment $\Gamma_{\chi, j}$ to an authenticated Merkle leaf handle $L_{\chi, j}$. Second, for each intermediate vector v , it links a SNARK-side vector commitment Γ_v^{vec} to the pre-query vector commitment μ_v . In our artifact, these equal-opening claims are batched into one QA-NIZK proof.

2.6. Freivalds' Algorithm

Freivalds' algorithm [1] is a classical randomized algorithm for verifying matrix multiplication. Given $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{k \times k}$, it samples $\mathbf{r} \leftarrow \mathbb{F}^k$ and checks whether $\mathbf{r}(\mathbf{AB}) = \mathbf{rC}$. The correctness relies on the following:

Lemma 2.2 (Schwartz–Zippel [6], [7]). *Let $P(X_1, \dots, X_m)$ be a non-zero polynomial of total degree d over a finite field $\mathbb{F}$. Then for $r_1, \dots, r_m \leftarrow \mathbb{F}$ chosen uniformly and independently,*

$$\Pr[P(r_1, \dots, r_m) = 0] \leq \frac{d}{|\mathbb{F}|}.$$

In our protocol, we use a structured variant: instead of sampling $\mathbf{r} \leftarrow \mathbb{F}^k$ uniformly, we sample one scalar $r \leftarrow \mathbb{F}$ and set $\mathbf{r} = (1, r, r^2, \dots, r^{k-1})$. This keeps the Freivalds challenge compact in the transcript and in the SNARK statement: the verifier records one field element rather than k independent ones. If $\mathbf{D} = \mathbf{AB} - \mathbf{C} \neq \mathbf{0}$, then for any non-zero column $\mathbf{d}_j$ of $\mathbf{D}$, the value $\langle \mathbf{r}, \mathbf{d}_j \rangle$ is a non-zero univariate polynomial in r of degree at most $k - 1$. By Lemma 2.2, it vanishes with probability at most $(k - 1)/|\mathbb{F}|$, which is negligible over the 256-bit fields used in our experiments.

3. Technical Overview

We provide a high-level description of our protocol for verifiable matrix multiplication before presenting the formal construction in Section 4. The central challenge is to verify that a claimed product $\mathbf{C} = \mathbf{AB}$ for matrices $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{k \times k}$ is correct, while keeping the verifier's

work substantially below the $\mathcal{O}(k^3)$ cost of recomputing the multiplication. Our approach combines three classical ideas—*randomized verification*, *error-correcting codes*, and *succinct commitment schemes*—in a way that decomposes the verification into components that are each efficient to check inside a SNARK circuit.

3.1. Starting Point: Freivalds’ Algorithm and Its Limitations

The classical starting point for verifying matrix multiplication is Freivalds’ algorithm [1]. Given matrices $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{k \times k}$ and a uniformly random vector $\mathbf{r} \in \mathbb{F}^k$, one checks whether $\mathbf{rAB} = \mathbf{rC}$. If $\mathbf{AB} \neq \mathbf{C}$, then $\mathbf{D} := \mathbf{AB} - \mathbf{C} \neq \mathbf{0}$, and the check $\mathbf{rD} = \mathbf{0}$ fails with probability at least $1 - k/|\mathbb{F}|$ by the Schwartz–Zippel lemma. This reduces a matrix-matrix verification to a vector-matrix computation, collapsing the problem from $\mathcal{O}(k^3)$ to $\mathcal{O}(k^2)$.

However, even the reduced $\mathcal{O}(k^2)$ cost is prohibitive when the check must be performed *inside a SNARK circuit*. In an RICS or Plonk arithmetization, each field multiplication corresponds to a constraint, so encoding the vector-matrix products $\mathbf{rA}$, $\mathbf{rB}$, and $\mathbf{rC}$ directly would still produce $\mathcal{O}(k^2)$ constraints. For matrices arising in modern deep-learning inference (where k can reach thousands), this SNARK circuit cost remains impractical.

3.2. From Vectors to Codewords: Leveraging Linear Codes

Our key observation is that the *linearity* of error-correcting codes can be exploited to push the bulk of the computation outside the circuit while retaining verifiability. Concretely, suppose we encode each row of a matrix $\mathbf{M} \in \mathbb{F}^{k \times k}$ using a systematic linear code $\mathcal{C} : \mathbb{F}^k \rightarrow \mathbb{F}^n$ with relative minimum distance δ . Write $\text{Enc}(\mathbf{M})$ for the matrix whose i -th row is $\mathcal{C}(\mathbf{M}_{i,*})$. Because $\mathcal{C}$ is $\mathbb{F}$ -linear, taking a linear combination of codewords yields the codeword of the corresponding linear combination of messages:

$$\sum_i r_i \cdot \mathcal{C}(\mathbf{M}_{i,*}) = \mathcal{C}\left(\sum_i r_i \cdot \mathbf{M}_{i,*}\right). \quad (2)$$

In other words, the “fold” operation $\text{Fold}(\text{Enc}(\mathbf{M}), \mathbf{r}) := \sum_i r_i \cdot \text{Enc}(\mathbf{M})_{i,*}$ equals $\text{Enc}(\mathbf{rM})$ exactly. This identity is the engine that allows us to *check vector-matrix products column-by-column in the encoded domain* without materializing the full vectors.

Commitment via Merkle trees. Each encoded matrix $\text{Enc}(\mathbf{M})$ is committed column-by-column using a Merkle tree:

$$\text{cm}_M \leftarrow \text{M.CM}(\text{Enc}(\mathbf{M})).$$

The prover sends input roots $(\text{cm}_A, \text{cm}_B, \text{cm}_C)$ before the challenge is sampled; these roots bind the encoded computation inputs and output using $\mathcal{O}(1)$ hash values.

3.3. Proximity-Based Verification via Random Sampling

After the Freivalds reduction, the verification task becomes: given committed encoded matrices $\text{Enc}(\mathbf{A}), \text{Enc}(\mathbf{B}), \text{Enc}(\mathbf{C})$ and a random challenge r , check that the intermediate vectors

$$\mathbf{x} = \mathbf{rA}, \quad \mathbf{y} = \mathbf{xB}, \quad \mathbf{z} = \mathbf{rC}$$

satisfy $\mathbf{y} = \mathbf{z}$. We now explain how to verify these relations in *sub-linear* time using the encoding.

The prover computes $\mathbf{x}, \mathbf{y}, \mathbf{z}$ outside the SNARK, encodes them, commits to the vectors and to their encoded leaves, and sends these commitments before the query tuple is sampled. The pre-query vector commitments are needed because the sampled Merkle openings alone only fix the queried coordinates; without a pre-query commitment to $\mathbf{x}, \mathbf{y}, \mathbf{z}$, the prover could choose vectors after seeing I that fit only those sampled positions. The verifier then samples an ordered tuple $I = (j_1, \dots, j_t) \leftarrow \$ [n]^t$ of column indices independently with replacement. For each sampled $j \in I$, the prover supplies the j -th column or coordinate as Groth16 witness data, publishes compact SNARK-side commitments bound to the relevant witnesses by the CC-SNARK backend, and provides auxiliary opening/linking material. The auxiliary material links full vector witnesses to the pre-query vector commitments and links sampled witnesses to authenticated external leaf handles.

At each queried position j , the verifier checks four *local consistency equations*:

- (i) $\text{Enc}(\mathbf{x})_j = \text{Fold}(\text{Enc}(\mathbf{A})^{(j)}, \mathbf{r})$, ensuring $\mathbf{x} = \mathbf{rA}$;
- (ii) $\text{Enc}(\mathbf{y})_j = \text{Fold}(\text{Enc}(\mathbf{B})^{(j)}, \mathbf{x})$, ensuring $\mathbf{y} = \mathbf{xB}$;
- (iii) $\text{Enc}(\mathbf{z})_j = \text{Fold}(\text{Enc}(\mathbf{C})^{(j)}, \mathbf{r})$, ensuring $\mathbf{z} = \mathbf{rC}$;
- (iv) $\text{Enc}(\mathbf{y})_j = \text{Enc}(\mathbf{z})_j$, ensuring the Freivalds check $\mathbf{y} = \mathbf{z}$.

The crucial point is that each of these checks involves only a *single column* of the encoded matrices—an inner product of length k (for the fold operations) plus a scalar comparison. Hence the total verifier work per query is $\mathcal{O}(k)$, and for t queries the total is $\mathcal{O}(tk)$.

Why does sampling suffice? The intermediate vectors are fixed before the query tuple is sampled. The vector-level CP-Link checks bind the vectors used by the SNARK to the pre-query vector commitments. At each sampled position, the SNARK relation recomputes $\text{Enc}(\mathbf{x})_j$, $\text{Enc}(\mathbf{y})_j$, and $\text{Enc}(\mathbf{z})_j$ from those vector witnesses, and checks that these values match the opened encoded leaves through the sampled linking checks. The verifier does not encode the full vectors; it verifies the SNARK proof and the sampled opening/linking material. If any of the underlying vector relations fails—say $\mathbf{y} \neq \mathbf{z}$ —then the codewords $\text{Enc}(\mathbf{y})$ and $\text{Enc}(\mathbf{z})$ are distinct. By the minimum-distance property of $\mathcal{C}$, they disagree in at least a δ -fraction of positions. A uniformly random query lands on a disagreement with probability $\geq \delta$, so t independent queries detect the inconsistency except with probability at most $(1 - \delta)^t$. By a union bound over the four checked relations, the overall

proximity-check failure probability is at most $4(1 - \delta)^t$. Setting $t = \lceil (\lambda + 2)/\log_2(1/(1 - \delta)) \rceil$ drives this union-bounded proximity term below $2^{-\lambda}$.

3.4. Putting It Together: The Full Protocol Pipeline

Combining the ingredients above, our protocol LAMP proceeds as follows (see Figure 1 for the formal description).

Commit inputs. The prover row-wise encodes $\mathbf{A}, \mathbf{B}, \mathbf{C}$ and sends roots $(\text{cm}_A, \text{cm}_B, \text{cm}_C)$. The roots are fixed before the Freivalds challenge and bind to those row-wise encodings under the linear code $\mathcal{C}$.

Challenge. The verifier samples a random field element $r \in \mathbb{F}$ and sends it to the prover. This induces the structured challenge vector $\mathbf{r} = (1, r, r^2, \dots, r^{k-1})$.

Respond. The prover computes the intermediate vectors $\mathbf{x} = \mathbf{r}\mathbf{A}$, $\mathbf{y} = \mathbf{x}\mathbf{B}$, $\mathbf{z} = \mathbf{r}\mathbf{C}$, commits to the vectors, encodes them, and commits to the encoded leaves $(\text{cm}_x, \text{cm}_y, \text{cm}_z)$.

Query. The verifier samples $I = (j_1, \dots, j_t) \leftarrow \$[n]^t$ and sends the sampled indices to the prover. The prover supplies the corresponding sampled columns or vector leaves as Groth16 witness data, together with public SNARK-side commitments and auxiliary Merkle openings for compact leaf handles.

SNARK verification. The verifier checks sampled RS evaluations for intermediate vectors and the four local consistency equations at each queried index through the Groth16 proof. The CC-SNARK interface binds the full vector witnesses and sampled witness values to public SNARK-side commitments. Merkle path validity, vector-level linking to the pre-query vector commitments, and sampled-linking consistency over the public leaf handles are auxiliary verifier checks bound to the same transcript. The prover then sends a succinct proof π plus the sampled opening/linking material, and the verifier accepts only if every component verifies.

Cost summary. The *SNARK circuit size*—the key metric governing prover cost—is $\mathcal{O}(t \cdot k)$ for sampled folds and RS evaluations. Vector-level CP-Link, Merkle authentication, and sampled-linking verification add $\mathcal{O}(t \log n + E_{\text{vec}}(k) + E_{\text{link}}(k))$ auxiliary online work. For the sampled-opening backend evaluated in this paper, the in-circuit checker is linear in k for fixed t and avoids the k^2 or k^3 circuit costs. The prover additionally performs the native matrix multiplication ($\mathcal{O}(k^3)$) and encoding ($\mathcal{O}(kn)$) outside the SNARK. The verifier’s core SNARK work is $\mathcal{O}(1)$ for pairing-based backends, or $\mathcal{O}(|x|)$ for transparent ones; total verification additionally depends on the selected commitment/linking backend. A detailed cost breakdown appears in Table 3 (Section 4.3).

4. Our Construction

In this section we present the formal construction of our verifiable matrix multiplication protocol LAMP. We begin by defining the relation chain that reduces the original matrix multiplication statement to a succinct SNARK-friendly relation (Section 4.1), then give the full protocol

Symbol	Description
$\mathcal{M}, \mathcal{W}$	Matrix objects $\{\mathbf{A}, \mathbf{B}, \mathbf{C}\}$ and intermediate vectors $\{\mathbf{x}, \mathbf{y}, \mathbf{z}\}$
$\mathcal{S}$	Sampled-object set $\mathcal{M} \cup \mathcal{W}$
$I = (j_1, \dots, j_t)$	Ordered tuple of sampled positions
cm_χ	Merkle root for committed object $\chi \in \mathcal{S}$
μ_v, ω_v	Pre-query commitment to intermediate vector v and its opening data
Γ_v^{vec}	SNARK-side commitment to the full intermediate vector v
$a_{\chi,j}$	Sampled object value used by the SNARK at position j
$L_{\chi,j}, \pi_{\chi,j}$	Authenticated leaf handle and opening path for object χ at position j
$\Gamma_{\chi,j}, \rho_{\chi,j}$	SNARK-side witness commitment to $a_{\chi,j}$ and its opening randomness
σ_I	Auxiliary proof for sampled opening/linking checks
$\mathbf{x}, \mathbf{w}$	Public instance and private witness for the SNARK relation
$E_{\text{vec}}(k), E_{\text{link}}(k)$	Costs for vector-level and sampled linking checks

TABLE 1. PROTOCOL-SPECIFIC NOTATION USED IN SECTION 4.

description (Section 4.2), and conclude with a complexity analysis (Section 4.3). Table 1 lists only the protocol-specific notation introduced in this section; basic field, matrix, code, Merkle, and SNARK notation is fixed in Section 2.

4.1. Relation Reduction Chain

Our construction proceeds through two successive reductions. Each reduction preserves completeness exactly and introduces a bounded soundness error that we analyze in Section 5.

Input-commit phase. The protocol starts by committing to the input and output matrices. The prover row-wise encodes $\mathbf{A}, \mathbf{B}, \mathbf{C}$ and publishes Merkle roots $\text{cm}_A, \text{cm}_B, \text{cm}_C$ before the Freivalds challenge is sampled. The input-root binding condition is that each root binds to $\text{Enc}(\mathbf{A}), \text{Enc}(\mathbf{B}), \text{Enc}(\mathbf{C})$ for fixed matrices $\mathbf{A}, \mathbf{B}, \mathbf{C} \in \mathbb{F}^{k \times k}$. This is the same committed-oracle viewpoint used in code-based proof systems such as Brakedown [5] and Orion [8]: once a root is fixed, any accepted opening must be the corresponding coordinate or column of the committed row-wise codeword. The multiplication checker then proves the relation among these input roots without arithmetizing the full matrix product. Our artifact reports the online checking cost after the roots are available; the raw cost of encoding matrices and building their roots is reported separately when measured and is not included in the online proof-byte claims. Throughout this section, $\text{MT.Com}(\text{Enc}(\mathbf{M}))$ is shorthand for the Merkle root over compact leaf handles $(L_{M,1}, \dots, L_{M,n})$, where each handle is binding to the corresponding encoded column $\text{Enc}(\mathbf{M})^{(j)}$. For an encoded vector, $\text{MT.Com}(\text{Enc}(v))$ analogously commits to handles for the length- n sequence of encoded coordinates. The sampled values themselves are opened only to the Groth16 relation as witness data; the auxiliary verifier sees handles and authentication paths.

Definition 4.1 (Sampled-opening backend). The online checker uses a black-box backend for sampled openings. For each intermediate vector $v \in \mathcal{W}$, the Groth16 relation

Contract item	What the main proof assumes
Input roots	$\text{cm}_A, \text{cm}_B, \text{cm}_C$ bind to fixed row-wise codewords before r .
Intermediate roots	μ_x, μ_y, μ_z and $\text{cm}_x, \text{cm}_y, \text{cm}_z$ are fixed before I .
Vector links	Accepted CP-Link checks bind each SNARK-side vector commitment Γ_v^{vec} to μ_v .
SNARK-side binding	Each sampled value used by Groth16 is bound to a public SNARK-side commitment.
Opening/linking	Accepted auxiliary checks bind each SNARK-side commitment to the opened leaf under the pre-query root.
Concrete backend	The artifact uses one batched linking proof with $\epsilon_{\text{link}} \leq \epsilon_{\text{ped}} + \epsilon_{\text{QA}} + Q/ \mathbb{F} $ for $Q = S t + \mathcal{W} $.

TABLE 2. BLACK-BOX BACKEND CONTRACT USED BY THE CONSTRUCTION AND SECURITY PROOF.

uses a private vector witness v and the CC-SNARK backend publishes a SNARK-side vector commitment Γ_v^{vec} to that vector. For each sampled object χ and query j , the Groth16 relation uses a private value $a_{\chi,j}$ and the CC-SNARK backend publishes a SNARK-side commitment $\Gamma_{\chi,j}$ to that value. The auxiliary verifier sees the committed root cm_χ , a leaf handle $L_{\chi,j}$, an opening path $\pi_{\chi,j}$, and an aggregate auxiliary proof σ_I . We write AuxV for the combined auxiliary verifier and LinkV for its CP-Link check after path verification. The backend must provide four guarantees:

SNARK-side binding. If the Groth16 verifier accepts, each intermediate vector witness v is bound to Γ_v^{vec} , and every sampled value used in the fold equations is bound to its public commitment $\Gamma_{\chi,j}$. This property is part of the selected SNARK backend and is charged to ϵ_{SNARK} .

Pre-query vector binding. The commitments μ_x, μ_y, μ_z are binding except with probability $\epsilon_{\text{vec}}(\lambda)$.

Position opening. If the auxiliary verifier accepts, each $L_{\chi,j}$ is authenticated as the leaf at position j under the pre-query root cm_χ .

CP-Link consistency. If both verifiers accept, then each value bound by Γ_v^{vec} is the same vector opened by μ_v , and each value bound by $\Gamma_{\chi,j}$ is the same value authenticated by $L_{\chi,j}$, except with probability $\epsilon_{\text{link}}(\lambda)$ over the backend security experiment. The theorem charges $\epsilon_{\text{cc}} \leq \epsilon_{\text{vec}} + \epsilon_{\text{link}}$.

We do not require a new global code-commitment primitive. Instead, **LAMP** is stated relative to the committed-codeword viewpoint of Brakedown [5] and Orion [8], plus the sampled-opening interface above. The concrete artifact instantiates the interface with pre-query vector commitments, SNARK-side vector commitments, Merkle openings, and a batched auxiliary linking proof; Appendix B gives the algebraic relation, batching argument, and online-byte accounting. The main construction and theorem use only the contract in Table 2.

Artifact backend in one paragraph. For the implementation evaluated in Section 6, each sampled SNARK-side commitment, each SNARK-side vector commitment, and each authenticated leaf handle is a Pedersen-style algebraic commitment to the same object as its counterpart. The auxiliary proof is one batched quasi-adaptive NIZK (QA-NIZK) proof for a linear relation [9] linking all $Q = |S|t + |\mathcal{W}|$ vector-level and sampled commitment pairs. Assuming CC-SNARK witness binding, binding of the

handle parameters, soundness of the QA-NIZK for the linear batch relation, and a domain-separated batch challenge sampled after the batch statement is fixed, the linking error is $\epsilon_{\text{link}} \leq \epsilon_{\text{ped}} + \epsilon_{\text{QA}} + Q/|\mathbb{F}|$. The CRS, handle generators, and proving/verifying keys are setup material; the reported byte counts in this paper are online proof bytes and exclude those serialized setup objects.

QA-NIZK, CRS, and generator model. The QA-NIZK is used only for the linear batch relation above, not for the matrix multiplication relation itself. Its CRS fixes the linear-relation verification parameters and the handle generator description used by the CP-Link backend. Those generators may be serialized directly or expanded deterministically from a public seed, depending on the concrete deployment; either way, the corresponding bytes are setup/key/CRS material, not online proof bytes. The current artifact reports online proof material and setup time, but not the serialized size of these setup objects.

Backend boundary. The security theorem is compositional: it proves the multiplication-checking protocol relative to any backend satisfying Definition 4.1. The input-commit phase may be run inside the same application protocol or by an upstream data pipeline, but the roots must be fixed before r is derived. For a $k \times k$ matrix with codeword length n , producing an input root costs $\Omega(kn)$ field work to encode rows and $\Omega(n)$ handle/hash work to form the Merkle tree. Static model-weight roots can amortize this cost over many inference requests; per-request activation and output roots cannot. The measurements in Section 6 therefore report the online multiplication-checking layer separately from raw-matrix encoding and root construction.

The original relation $\mathcal{R}_{\text{orig}}$. The starting point is the statement that the prover knows matrices whose *encoded* commitments satisfy the multiplication relation. The public instance is $(\text{cm}_A, \text{cm}_B, \text{cm}_C)$ and the witness is (A, B, C) ; the relation $\mathcal{R}_{\text{orig}}$ requires:

$$\begin{aligned} \text{cm}_A &= \text{MT.Com}(\text{Enc}(A)), & \text{cm}_B &= \text{MT.Com}(\text{Enc}(B)), \\ \text{cm}_C &= \text{MT.Com}(\text{Enc}(C)), & \mathbf{AB} &= \mathbf{C}. \end{aligned}$$

Thus a matrix commitment binds the verifier to the row-wise codeword matrix rather than to the raw matrix alone. Since $\mathcal{C}$ is systematic, the raw matrix is also determined by the first k coordinates of each committed row. Proving membership in $\mathcal{R}_{\text{orig}}$ directly inside a SNARK would require $\mathcal{O}(k^3)$ multiplication constraints, which is exactly what we aim to avoid.

Reduction 1: Freivalds’ randomization ($\mathcal{R}_{\text{orig}} \rightsquigarrow \mathcal{R}_{\text{Freivalds}}$). Using the verifier’s random challenge $r \in \mathbb{F}$, define $\mathbf{r} = (1, r, r^2, \dots, r^{k-1})$ and introduce three intermediate vectors

$$\mathbf{x} = \mathbf{rA}, \quad \mathbf{y} = \mathbf{xB}, \quad \mathbf{z} = \mathbf{rC}.$$

The Freivalds relation checks

$$\mathbf{rA} = \mathbf{x}, \quad \mathbf{xB} = \mathbf{y}, \quad \mathbf{rC} = \mathbf{z}, \quad \mathbf{y} = \mathbf{z}.$$

By the Schwartz–Zippel lemma, if $\mathbf{AB} \neq \mathbf{C}$, then the final equality fails except with probability at most $(k-1)/|\mathbb{F}|$

over the choice of r . This removes the cubic matrix-product check, but checking the three vector-matrix products directly would still require quadratic circuit work. No commitment to $\mathbf{x}, \mathbf{y}, \mathbf{z}$ is needed at this algebraic step; those commitments are introduced next only to make the sampled proximity checks non-adaptive.

Reduction 2: Proximity lifting ($\mathcal{R}_{\text{Freivalds}} \rightsquigarrow \mathcal{R}_{\text{Prox}}$). We now encode the matrices and intermediate vectors using a systematic $\mathbb{F}$ -linear code $\mathcal{C} : \mathbb{F}^k \rightarrow \mathbb{F}^n$ with relative distance δ , and replace the global vector-matrix checks with *local* consistency checks at random positions. Let $I = (j_1, \dots, j_t) \leftarrow \$_{[n]}^t$ be the ordered tuple of queried column indices sampled independently with replacement, and let $\mathcal{M} = \{\mathbf{A}, \mathbf{B}, \mathbf{C}\}$ denote the matrix set and $\mathcal{W} = \{\mathbf{x}, \mathbf{y}, \mathbf{z}\}$ the intermediate-vector set. Let $\mathcal{S} = \mathcal{M} \cup \mathcal{W}$ be the set of objects opened at the sampled positions. The matrix roots $\text{cm}_A, \text{cm}_B, \text{cm}_C$ are fixed before r is sampled. After seeing r but before seeing I , the prover commits to $\mathbf{x}, \mathbf{y}, \mathbf{z}$ by sending pre-query vector commitments μ_x, μ_y, μ_z and encoded-vector roots $\text{cm}_x, \text{cm}_y, \text{cm}_z$. The roots $\text{cm}_x, \text{cm}_y, \text{cm}_z$ fix the encoded leaves that will later be opened, but by themselves they do not force the Groth16 witness vectors $\mathbf{x}, \mathbf{y}, \mathbf{z}$ to be chosen before I . The vector-level CP-Link checks close this gap: the CC-SNARK backend commits to the full vector witnesses used by Groth16, and the auxiliary proof links those SNARK-side vector commitments to the pre-query commitments μ_x, μ_y, μ_z . Without this vector-level link, a malicious prover could wait for I and choose vectors that interpolate only the sampled coordinates.

The *public instance* is

$$\mathbf{x} = (\{\text{cm}_\chi\}_{\chi \in \mathcal{S}}, \mu_x, \mu_y, \mu_z, r, I, \{\Gamma_v^{\text{vec}}\}_{v \in \mathcal{W}}, \{\Gamma_{\chi,j}\}_{\chi \in \mathcal{S}, j \in I}).$$

and the private arithmetic-relation witness is

$$\mathbf{w} = (\mathbf{x}, \mathbf{y}, \mathbf{z}, \{a_{\chi,j}\}_{j \in I, \chi \in \mathcal{S}}).$$

The auxiliary opening/linking proof data is

$$\text{aux}_I = (\{L_{\chi,j}, \pi_{\chi,j}\}_{j \in I, \chi \in \mathcal{S}}, \sigma_I).$$

For a matrix commitment, $a_{\chi,j} \in \mathbb{F}^k$ is the sampled encoded column $\text{Enc}(\chi)^{(j)}$ used by the fold equation. For a vector commitment, $a_{\chi,j} \in \mathbb{F}$ is the sampled encoded coordinate. The opening randomness $\rho_v^{\text{vec}}, \rho_{\chi,j}$, and ω_v belong to the SNARK-side or external commitments and to the auxiliary linking witness; they are not RICS witness variables checked by in-circuit opening constraints. The resulting proximity-checking relation, denoted $\mathcal{R}_{\text{Prox}}$, requires $(\mathbf{x}, \mathbf{w}, \text{aux}_I)$ to satisfy all of the following clauses. In our implementation, clauses (RS)–(EQ) are checked by Groth16; clauses (PATH) and (LINK) are checked by AuxV over the

same transcript and SNARK-side commitments.

$$\begin{aligned} \forall v \in \mathcal{W}, j \in I: & \quad a_{v,j} = \text{Enc}(v)_j, & (\text{RS}) \\ \forall j \in I: & \quad a_{x,j} = \text{Fold}(a_{A,j}, \mathbf{r}), & (\text{FOLD-A}) \\ \forall j \in I: & \quad a_{y,j} = \text{Fold}(a_{B,j}, \mathbf{x}), & (\text{FOLD-B}) \\ \forall j \in I: & \quad a_{z,j} = \text{Fold}(a_{C,j}, \mathbf{r}), & (\text{FOLD-C}) \\ \forall j \in I: & \quad a_{y,j} = a_{z,j}, & (\text{EQ}) \\ \forall \chi \in \mathcal{S}, j \in I: & \quad \text{MT.Open}(\pi_{\chi,j}, L_{\chi,j}, \text{cm}_\chi, j) = 1 & (\text{PATH}) \\ \text{LinkV}(\mathbf{x}, I, ((\Gamma_v^{\text{vec}}, \mu_v))_{v \in \mathcal{W}}, ((\Gamma_{\chi,j}, L_{\chi,j}))_{\chi \in \mathcal{S}, 1 \leq j \leq t}, \sigma_I) = 1 & (\text{LINK}) \end{aligned}$$

Here $\forall j \in I$ ranges over every sampled position in the ordered tuple, including repetitions. The CC-SNARK backend is the bridge between the private Groth16 witness and the auxiliary opening world. It binds each full intermediate vector v to Γ_v^{vec} , and each sampled value $a_{\chi,j}$ to $\Gamma_{\chi,j}$. Equation (PATH) uses the selected position-opening verifier on the compact leaf handle $L_{\chi,j}$. For the transparent Merkle instantiation, this is MT.Verify with the queried index included in the leaf domain. Equation (LINK) is the backend check that links the SNARK-side vector commitments Γ_v^{vec} to the pre-query vector commitments μ_v , and links each sampled SNARK-side commitment $\Gamma_{\chi,j}$ to the authenticated leaf handle $L_{\chi,j}$. This vector-level link is what ties the vectors used by the circuit to the pre-query vector commitments fixed before I was sampled. Sending μ_x, μ_y, μ_z alone would not force the circuit witness to use the committed vectors. Equation (RS) computes the sampled RS/code evaluation for each intermediate vector inside the circuit. In the LAMP_{QA} artifact, sampled RS evaluations and fold equations are included in the Groth16 relation, while vector-level CP-Link checks, Merkle path verification, and sampled linking are auxiliary checks bound to the same transcript and to the SNARK-side commitments. The security proof uses only vector-level link consistency, sampled witness-to-leaf link consistency, and position binding of the opening layer.

The critical property is that the Groth16 relation involves only $\mathcal{O}(|I| \cdot k)$ constraints for the fold operations, sampled RS evaluations, and equality checks. Merkle authentication, vector-level CP-Link, and sampled-linking verification are auxiliary online checks whose cost is reported separately. Since $|I|$ is a security-parameter-sized constant independent of k , the in-circuit checking relation is *linear in k* .

Intermediate-binding cost. The vector-level links in (LINK) force the vectors used by Groth16 to be the same vectors committed by μ_x, μ_y, μ_z before the query tuple is sampled. The CC-SNARK witness binding of $\Gamma_{\chi,j}$ and the sampled part of (LINK) then force each sampled external opening to equal the value used by the in-circuit fold relation; for intermediate vectors, (RS) also forces this value to be the corresponding code evaluation $\text{Enc}(v)_j$. Without the pre-query vector commitments, a cheating prover could choose sampled openings first and then interpolate a vector that matches only those positions, defeating the distance-based soundness argument. Analogously, the input matrix roots $\text{cm}_A, \text{cm}_B, \text{cm}_C$ must be fixed commitments to row-wise encodings of $\mathbf{A}, \mathbf{B}, \mathbf{C}$ before the Freivalds challenge. For Reed–Solomon-style encodings, a sampled coordinate $\text{Enc}(v)_j$ can be computed by evaluating the degree- $(k-1)$

interpolation polynomial at the queried domain point, costing $\mathcal{O}(k)$ constraints per sampled coordinate. Thus the sampled intermediate checks contribute $\mathcal{O}(tk)$ constraints, while the vector-level links are auxiliary backend checks. A fallback backend that proves full dense codeword validity inside the circuit would instead require $\mathcal{O}(k(n-k))$ work for a dense generator matrix. We discuss the trade-off in Section 4.3.

4.2. Protocol Description

We now give the full specification of LAMP. Let $\mathcal{C} : \mathbb{F}^k \rightarrow \mathbb{F}^n$ be a systematic linear code with relative distance δ and MT.Com denote the Merkle tree commitment. Let $\mathcal{M} = \{\mathbf{A}, \mathbf{B}, \mathbf{C}\}$ be the matrix set, $\mathcal{W} = \{\mathbf{x}, \mathbf{y}, \mathbf{z}\}$ the intermediate-vector set, and $\mathcal{S} = \mathcal{M} \cup \mathcal{W}$ the objects opened at sampled positions. The prover first commits to $\text{Enc}(\mathbf{A}), \text{Enc}(\mathbf{B}), \text{Enc}(\mathbf{C})$ and sends the input roots $(\text{cm}_A, \text{cm}_B, \text{cm}_C)$ before the challenge is derived. The protocol LAMP operates between a prover $\mathcal{P}$ and verifier $\mathcal{V}$ as shown in Figure 1.

Remark on the SNARK relation. The SNARK proof π attests to the in-circuit part of $\mathcal{R}_{\text{prox}}$ defined in Section 4.1; full transcript acceptance additionally requires the auxiliary opening and linking verifiers to accept. The implemented verifier checks three groups of conditions, split between the Groth16 circuit and auxiliary verifiers. First, constraints (RS)–(EQ) verify the sampled RS, fold, and equality checks at each $j \in I$ inside the Groth16 circuit. The CC-SNARK backend binds the full vector witnesses to Γ_v^{vec} and binds each sampled value used by the circuit to the public commitment $\Gamma_{\chi,j}$ through the Groth16 verification equation; these bindings are not separate arithmetic constraints. Second, the auxiliary backend checks (LINK), linking each Γ_v^{vec} to μ_v and linking each public sampled commitment $\Gamma_{\chi,j}$ to the authenticated external leaf handle $L_{\chi,j}$. Third, the auxiliary opening verifier checks (PATH) for all Merkle opening proofs over the leaf handles. The public statement $\times$ consists of the commitment roots $\{\text{cm}_\chi\}_{\chi \in \mathcal{S}}$, the intermediate vector commitments μ_x, μ_y, μ_z , the challenge r , the query tuple I , the SNARK-side vector commitments $\{\Gamma_v^{\text{vec}}\}$, and the sampled SNARK-side commitments $\{\Gamma_{\chi,j}\}$. The relation witness w consists of the intermediate vectors and the sampled matrix columns or vector leaves used by the fold equations. The randomness $\{\rho_v^{\text{vec}}\}, \{\rho_{\chi,j}\}$, and $\{\omega_v\}$ is consumed by the CC-SNARK/linking backend, not by an in-circuit opening check.

Non-interactive variant. The protocol as described is a three-round public-coin interactive protocol. By applying the Fiat–Shamir heuristic [10] in the random oracle model, it can be made non-interactive using domain-separated transcript hashes:

$$\begin{aligned} \text{cm}_{ABC} &\leftarrow \text{Hash}(\text{pp}, \text{cm}_A, \text{cm}_B, \text{cm}_C), \\ r &\leftarrow \text{Hash}(\text{FMM-r}, \text{cm}_{ABC}). \end{aligned}$$

and

$$\begin{aligned} \text{cm}_{XYZ} &\leftarrow \text{Hash}(\text{pp}, r, \mu_x, \mu_y, \mu_z, \text{cm}_x, \text{cm}_y, \text{cm}_z), \\ T_{\text{qry}} &\leftarrow (\text{cm}_{ABC}, \text{cm}_{XYZ}), \\ I &\leftarrow \text{Sample}_t(\text{Hash}(\text{FMM-I}, T_{\text{qry}})). \end{aligned}$$

The resulting non-interactive proof consists of the commitment roots, the intermediate vector commitments, the transcript-derived challenges, the sampled SNARK-side commitments $\{\Gamma_{\chi,j}\}$, and the SNARK proof π . Merkle leaf handles, authentication paths, and sampled-linking material are auxiliary proof data verified by the corresponding backend. Raw sampled matrix columns are not serialized as proof data; they remain private Groth16 witness values bound to the public SNARK-side commitments.

Security of the non-interactive variant. A potential concern with the Fiat–Shamir transform for proximity-based protocols is a *grinding attack*: a malicious prover could try many transcript prefixes until the derived challenge r or query tuple I is favorable. The random-oracle theorem accounts for this by the standard $\epsilon_{\text{FS}}(q_H)$ loss in Theorem 5.1. For intuition, consider grinding only the intermediate commitments after r is fixed. Suppose the adversary has committed to $(\text{cm}_A, \text{cm}_B, \text{cm}_C)$ (and hence to fixed encoded matrices) and, after receiving r , is searching for intermediate vector commitments and encoded-vector roots such that the transcript-derived query tuple I avoids the disagreement set D of some fold or equality check. If any of the four checks fails globally, then $|D| \geq \delta n$ by the minimum-distance property. For a single random query $j \in [n]$, the probability of $j \notin D$ is at most $1 - \delta$; for t independent queries it is $(1 - \delta)^t$. Choosing t as in Section 5 makes the union-bounded proximity term $4(1 - \delta)^t$ at most $2^{-\lambda}$. Each trial of the grinding attack produces an independently random I in the random oracle model because the query hash binds the full transcript, so the probability that q grinding attempts yield a favorable I is at most $q \cdot 4(1 - \delta)^t$. For any PPT adversary, q is polynomially bounded in λ , giving a total success probability of $\text{negl}(\lambda)$. Hence the non-interactive protocol retains the soundness bound of Theorem 5.1 in the random oracle model, plus the stated Fiat–Shamir loss.

Zero-knowledge extension. The protocol as stated is a *verifiable computation* protocol and does not conceal the inputs $\mathbf{A}, \mathbf{B}, \mathbf{C}$ from the verifier. We do not claim zero knowledge for LAMP. A zero-knowledge variant would require replacing the Merkle tree commitments with hiding commitments, using a zero-knowledge SNARK backend for π , and masking the intermediate vectors and sampled openings so that the transcript does not reveal linear information about the committed matrices. The auxiliary opening/linking material would also need a separate simulation argument. We therefore treat zero knowledge as future work rather than as a property of the construction analyzed in this paper.

4.3. Complexity Analysis

We analyze the costs of LAMP.

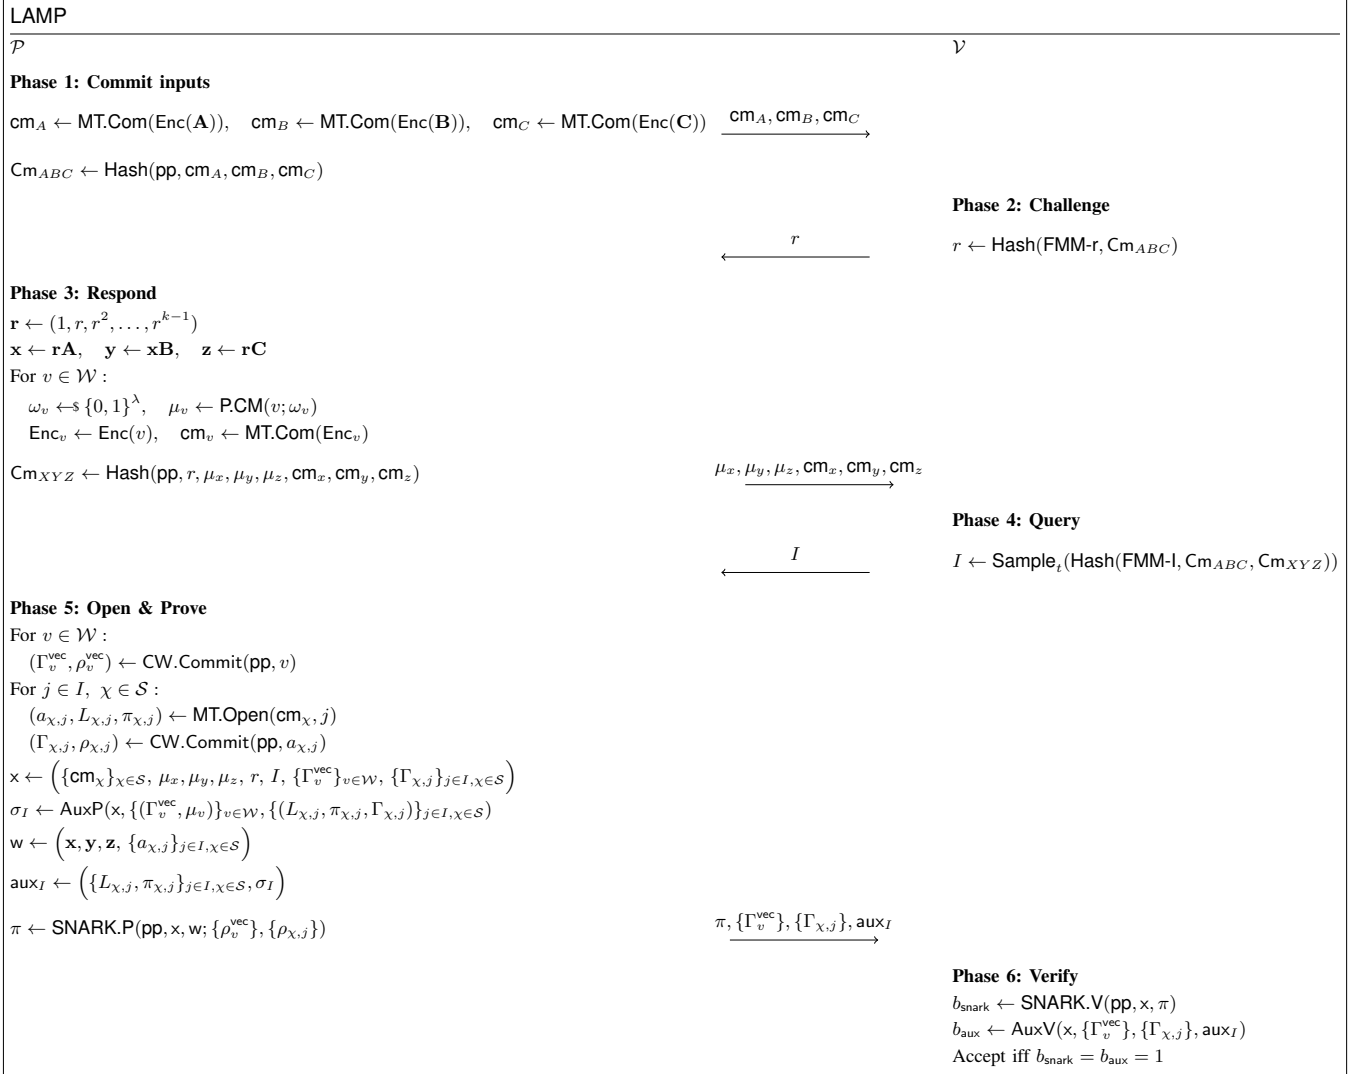


Figure 1. The LAMP protocol for verifiable matrix multiplication. $\mathcal{S} = \mathcal{M} \cup \mathcal{W}$ is the set of objects opened at sampled positions, $t = |I|$ is the query complexity, and $\mathbf{x}, \mathbf{w}$ denote the public statement and witness for the inner SNARK.

Prover computation. We separate the cost of producing the claimed output matrix and input roots from the cost of proving that the roots satisfy the multiplication claim. Computing $\mathbf{C} = \mathbf{AB}$ is application work shared by all approaches and is not part of the checking-layer comparison. Likewise, producing input roots for raw matrices is an input-commitment cost: with linear-time codes it takes $\mathcal{O}(kn)$ field work for a $k \times k$ matrix, plus handle and hash work to form the root. After those roots are fixed, the prover's dominant online checking-layer costs are:

- (i) *Vector computation:* Computing $\mathbf{x}, \mathbf{y}, \mathbf{z}$ costs $\mathcal{O}(k^2)$ in total.
- (ii) *Intermediate-vector roots:* Encoding $\mathbf{x}, \mathbf{y}, \mathbf{z}$ and forming their roots costs $\mathcal{O}(n)$ for constant-rate linear codes.
- (iii) *SNARK proving:* The circuit size is $\mathcal{O}(t \cdot k)$ for the fold checks and sampled RS evaluations. Vector-level CP-Link, Merkle authentication, CC-SNARK witness

commitments, and sampled-linking work are backend costs included in the reported proving, verification, and online-byte breakdowns.

Verifier computation. The verifier performs SNARK verification together with any commitment/linking checks required by the chosen backend. For a pairing-based SNARK whose public input consists only of the roots, challenge, and query indices, the core SNARK verification is $\mathcal{O}(1)$. For transparent or linking-heavy backends, verification may scale with the public input or linking proof bytes; the concrete behavior is reported in Section 6.

Communication. The communication depends on the commitment layer. With the sampled-opening backend evaluated in this paper, the proof contains the SNARK proof, vector commitments, public SNARK-side commitments, authentication paths, and auxiliary linking material for the queried objects. The raw sampled matrix columns contain $\mathcal{O}(tk)$ field elements, but they are Groth16 wit-

Component	Prover	Verifier
Input roots	$\mathcal{O}(kn)$ root generation	fixed before r
Intermediate vectors	$\mathcal{O}(k^2)$	—
Vector roots	$\mathcal{O}(n)$	—
SNARK circuit size	$\mathcal{O}(tk)$	—
Auxiliary opening/linking	$\mathcal{O}(t \log n + E_{\text{vec}}(k) + E_{\text{link}}(k))$	backend-dependent
Core SNARK verification	—	$\mathcal{O}(1)$
Online total	checking $\mathcal{O}(k^2 + n + tk + E_{\text{vec}} + E_{\text{link}})$	backend-dependent

TABLE 3. ASYMPTOTIC COSTS AFTER SEPARATING INPUT-ROOT GENERATION FROM ONLINE MULTIPLICATION CHECKING. THE ONLINE TOTAL EXCLUDES NATIVE COMPUTATION OF $\mathbf{C} = \mathbf{A}\mathbf{B}$ AND RAW MATRIX ENCODING/ROOT CONSTRUCTION. k : MATRIX DIMENSION; n : CODEWORD LENGTH; t : QUERY COUNT; E_{vec} AND E_{link} ARE BACKEND COSTS.

ness values rather than serialized proof bytes. The serialized auxiliary material is compact: SNARK-side commitments, leaf handles, Merkle paths, and link proofs scale as $\mathcal{O}(t \log n + E_{\text{link}}(k))$ for fixed security parameter. This accounting explains the sub-megabyte online proof sizes reported in Section 6 and Appendix C; making the sampled columns public would instead add $\Theta(tk)$ field elements.

Circuit size comparison. Table 3 summarizes the asymptotic costs. Naïve arithmetization uses $\mathcal{O}(k^3)$ constraints. Freivalds’ algorithm alone uses $\mathcal{O}(k^2)$. The online checker achieves

$$\mathcal{O}(tk) \text{ in-circuit work}$$

plus $\mathcal{O}(t \log n + E_{\text{vec}}(k) + E_{\text{link}}(k))$ auxiliary opening/linking work. The in-circuit term covers the fold checks and sampled RS evaluations. The auxiliary term covers vector-level CP-Link, Merkle authentication, and sampled-linking verification over compact sampled commitments and leaf handles. The older notation $E_{\text{cc}}(k)$ denotes the combined vector-level and sampled linking cost, with $E_{\text{cc}}(k) = E_{\text{vec}}(k) + E_{\text{link}}(k)$ for the current artifact. For dense general codes, a full in-circuit codeword-validity fallback gives $E_{\text{cc}}(k) = \mathcal{O}(k(n - k))$, which is $\mathcal{O}(k^2)$ at constant rate. For the sampled-opening backend evaluated in this paper, the online sampled encoding work is already included in the $\mathcal{O}(tk)$ term. Setting

$$t = \left\lceil \frac{\lambda + 2}{\log_2(1/(1 - \delta))} \right\rceil$$

then gives $\mathcal{O}(k)$ in-circuit work and $\mathcal{O}(\log k)$ serialized auxiliary opening material for fixed security parameter, with concrete sizes discussed in Section 6.

5. Security Analysis

We prove that the protocol LAMP (Figure 1) constitutes a secure interactive argument for the matrix multiplication relation $\mathcal{R}_{\text{orig}}$ over encoded input roots and the sampled-opening backend of Definition 4.1. Soundness and completeness are established in Theorem 5.1 below; knowledge

soundness is discussed separately in Remark 2. As discussed in Section 4.1, the input roots $\text{cm}_A, \text{cm}_B, \text{cm}_C$ bind to row-wise encodings of fixed matrices $\mathbf{A}, \mathbf{B}, \mathbf{C}$. This is the committed-input model used by code-based proof systems such as Brakedown [5] and Orion [8]; LAMP verifies the multiplication relation among these committed objects. For the intermediate vectors, the prover sends pre-query vector commitments to $\mathbf{x}, \mathbf{y}, \mathbf{z}$ before the query tuple is sampled. The CC-SNARK backend binds each vector witness used by Groth16 to a public SNARK-side vector commitment Γ_v^{vec} , and the auxiliary CP-Link proof links Γ_v^{vec} to the pre-query commitment μ_v . At each sampled coordinate, the CC-SNARK backend also binds the matrix column or vector leaf used by the circuit to a public SNARK-side commitment $\Gamma_{\chi,j}$. The auxiliary verifier authenticates a compact external Merkle leaf handle $L_{\chi,j}$ and checks that it is linked to the SNARK-side commitment. For intermediate vectors, the SNARK also computes $\text{Enc}(v)_j$ internally before the sampled value is linked to the authenticated handle.

Scope of the security claim. The theorem is a statement-soundness result for committed matrix multiplication, not a standalone argument-of-knowledge or zero-knowledge theorem. The public objects are the encoded input roots and the transcript-derived challenges; the committed objects are the row-wise encodings of $\mathbf{A}, \mathbf{B}, \mathbf{C}$ and the intermediate encoded vectors. The backend supplies vector-level links to the pre-query commitments and sampled consistency between the SNARK-side commitments and the external Merkle openings. Witness extraction and hiding require the additional composition assumptions described in Remark 2 and are outside the main theorem. In short, $\text{cm}_A, \text{cm}_B, \text{cm}_C$ are input roots fixed before r ; μ_x, μ_y, μ_z and $\text{cm}_x, \text{cm}_y, \text{cm}_z$ are fixed before I ; the full intermediate vectors and sampled raw columns are private Groth16 witness values linked to pre-query commitments and authenticated leaf handles.

Soundness-critical ordering. The proof uses the following timing invariants. First, the input roots $\text{cm}_A, \text{cm}_B, \text{cm}_C$ are fixed before the Freivalds challenge r is sampled. Second, the intermediate vector commitments and encoded-vector roots are fixed before the query tuple $I = (j_1, \dots, j_t) \leftarrow \$ [n]^t$ is sampled; conditioned on accepting vector-level CP-Link proofs, the vectors used by Groth16 are the vectors fixed by μ_x, μ_y, μ_z independently of the sampled positions, and accepted auxiliary linking proofs bind sampled external openings to the corresponding SNARK-side commitments. For intermediate vectors, the vector-level links and in-circuit RS checks force the value behind each SNARK-side commitment to be the corresponding coordinate of $\text{Enc}(\mathbf{x}), \text{Enc}(\mathbf{y}), \text{Enc}(\mathbf{z})$. Third, the entries of I are sampled independently with replacement, as defined in Section 2.1. If an implementation samples without replacement, the proximity failure probability only decreases; the theorem states the simpler with-replacement bound.

Backend failure event. The backend bad event is the only non-standard composition point. After the transcript prefix fixes $(\text{Cm}_{ABC}, r, \text{Cm}_{XYZ}, I)$, the adversary must

output SNARK-side vector commitments, sampled SNARK-side commitments, Merkle openings, and an auxiliary linking proof for the vector-level and sampled objects. A backend failure occurs if all these checks accept while some value used in the Groth16 fold relation is not the value fixed by the corresponding pre-query vector commitment or authenticated at the corresponding leaf position. Definition 4.1 charges this event to ϵ_{cc} ; Appendix A expands the event for the concrete backend. For the artifact backend summarized in Section 4, the batched CP-Link component contributes $\epsilon_{link} \leq \epsilon_{ped} + \epsilon_{QA} + Q/|\mathbb{F}|$, where $Q = |\mathcal{S}|t + |\mathcal{W}|$ is the number of sampled SNARK-side/leaf commitment pairs plus vector-level commitment pairs. Here ϵ_{ped} is the binding error of the algebraic handle parameters, ϵ_{QA} is the soundness error of the batched QA-NIZK relation, and the $Q/|\mathbb{F}|$ term is the standard randomized-batch failure probability after the batch statement is fixed.

Theorem 5.1 (Security). *The protocol LAMP satisfies the following properties:*

- (i) **Perfect completeness.** *If $\mathbf{AB} = \mathbf{C}$, an honest prover always makes the verifier accept.*
- (ii) **Interactive computational soundness.** *For any PPT adversary $\mathcal{P}^*$ given public input roots for $\mathbf{A}, \mathbf{B}, \mathbf{C}$, let Acc be the event that the verifier accepts a false committed matrix statement. Then:*

$$\Pr[\text{Acc} \mid \mathbf{AB} \neq \mathbf{C}] \leq \epsilon_{\text{SNARK}}(\lambda) + \epsilon_{cc}(\lambda) + \epsilon_{\text{bind}}(\lambda) + \frac{k-1}{|\mathbb{F}|} + 4(1-\delta)^t,$$

Here $t = |I|$ is the query complexity and δ is the relative minimum distance of $\mathcal{C}$. The terms ϵ_{SNARK} , ϵ_{cc} , and ϵ_{bind} denote, respectively, the soundness error of the SNARK backend including CC-SNARK witness binding, the vector-level and sampled linking consistency error of the selected intermediate backend, and the union-bounded position-binding error of the opening layer over all roots and openings used in the protocol. In the notation of Definition 4.1, $\epsilon_{cc} \leq \epsilon_{\text{vec}} + \epsilon_{\text{link}}$. For the artifact backend, $\epsilon_{\text{link}} \leq \epsilon_{ped} + \epsilon_{QA} + (|\mathcal{S}|t + |\mathcal{W}|)/|\mathbb{F}|$; Appendix B gives the full batching equation.

- (iii) **Fiat-Shamir soundness.** *In the random oracle model, the non-interactive variant described in Section 4.2 satisfies the same bound with an additional Fiat-Shamir loss term $\epsilon_{\text{FS}}(q_H)$, where q_H is the adversary's number of random-oracle queries.*

Corollary 5.2 (Evaluated backend). *For the backend used in the artifact, Theorem 5.1 applies with*

$$\epsilon_{cc} \leq \epsilon_{\text{vec}} + \epsilon_{ped} + \epsilon_{QA} + (|\mathcal{S}|t + |\mathcal{W}|)/|\mathbb{F}|.$$

The bound assumes the pre-query vector commitment is binding, the algebraic handle parameters are binding, the QA-NIZK is sound for the stated linear batch relation, and the batch randomizer is sampled after the commitment-pair statement is fixed. In the non-interactive artifact, the Fiat-Shamir loss term $\epsilon_{\text{FS}}(q_H)$ covers both the protocol challenges (r, I) and the backend batch randomizer, using separate domain tags.

The reported scaling rows use $t = 128$; if $\delta = 1/2$, this gives proximity term 2^{-126} , while a conservative per-term 2^{-131} target uses $t = 133$. A strict deployment must also instantiate the SNARK, pre-query vector commitment, CP-Link backend, and hash parameters so their union-bounded errors fit the aggregate security budget.

Proof sketch. Completeness follows from linearity of the Freivalds reduction and of the encoding: honest openings satisfy every sampled fold and equality check. For soundness, condition away SNARK, vector-level linking, sampled linking, and position-binding failures. Then the input roots fix $\mathbf{A}, \mathbf{B}, \mathbf{C}$ before r , the intermediate commitments fix $\mathbf{x}, \mathbf{y}, \mathbf{z}$ before I , and accepted sampled openings are the same values used by the Groth16 fold relation. Schwartz-Zippel gives the Freivalds term $(k-1)/|\mathbb{F}|$; conditioned on Freivalds not failing, at least one of the four encoded relations is false, so t random coordinates miss its disagreement set with probability at most $(1-\delta)^t$. Union-bounding over the four relations gives the proximity term. Appendix A gives the detailed proof and the Fiat-Shamir random-oracle reduction.

Remark 1 (Fiat-Shamir soundness). The proof of the main bound in Theorem 5.1 is for the public-coin interactive protocol. The non-interactive variant in Section 4.2 derives both r and I from domain-separated hashes of the full transcript prefix. The artifact backend also derives its batch randomizer with a separate domain tag after the commitment-pair batch statement is fixed. In the random oracle model, the same proof applies with the standard Fiat-Shamir loss $\epsilon_{\text{FS}}(q_H)$, where q_H is the adversary's number of random-oracle queries. This term accounts for grinding over transcript prefixes and is negligible for polynomial q_H when the challenge and query hashes have λ -bit security.

Remark 2 (Knowledge Soundness). The theorem above is a statement-soundness result relative to input roots that bind to row-wise encoded matrices. A full *argument of knowledge* can be obtained by composing LAMP with an extractable input-commitment interface and a knowledge-sound SNARK. In that setting, the input-commitment extractor yields openings of $\text{cm}_A, \text{cm}_B, \text{cm}_C$ to row-wise encoded matrices, while the SNARK extractor yields the intermediate vectors and queried openings used in $\mathcal{R}_{\text{Prox}}$. The same soundness argument then implies that any prover accepted with non-negligible probability is bound, except with the stated error, to matrices satisfying $\mathbf{AB} = \mathbf{C}$. We treat this as an orthogonal composition issue and focus the main theorem on the multiplication-checking protocol.

6. Implementation and Evaluation

We implement the multiplication-checking layer of LAMP and evaluate the current artifact, denoted LAMP_{QA} , against an optimized SNARK-in-circuit Freivalds baseline [1]. LAMP_{QA} uses multi-column Merkle openings, a Groth16 circuit proof, and the sampled-opening backend from Section 4. All main experiments use a 256-bit prime field, rate $\rho = 1/2$, codeword length $n = 2k$, and $t = 128$

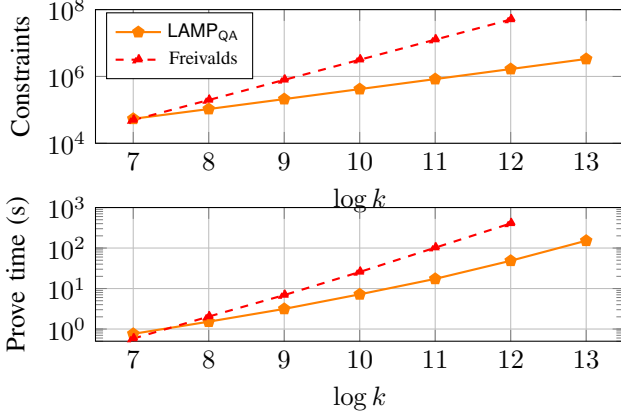


Figure 2. Primary matrix benchmark. With the sampled-opening backend, LAMP_{QA} follows the expected $\mathcal{O}(k)$ checking-layer growth, whereas the Freivalds SNARK baseline follows $\mathcal{O}(k^2)$ growth.

sampling positions. The experiments were run on a Linux server with an AMD EPYC 7B13 CPU (32 cores) and 251.69 GiB RAM.

These measurements include the online sampled-handle path described in Section 4: pre-query vector commitments fix $\mathbf{x}, \mathbf{y}, \mathbf{z}$ before I , vector-level CP-Link connects those commitments to the Groth16 vector witnesses, the CC-SNARK backend binds every sampled matrix column or vector leaf to a public SNARK-side commitment, Merkle openings authenticate compact external leaf handles, and sampled linking checks connect the sampled views. The raw sampled matrix columns are SNARK witness data, not serialized auxiliary proof bytes. All byte counts in this section are online proof bytes: they exclude raw-matrix encoding and input-root construction for $\mathbf{cm}_A, \mathbf{cm}_B, \mathbf{cm}_C$, proving/verifying keys, serialized CRS data, and generator material. They should not be read as complete end-to-end deployment costs from raw matrices. Appendix C records parameter and measurement coverage details.

Primary matrix benchmark. Figure 2 shows the main scaling behavior. Freivalds constraints grow by approximately $4\times$ with each doubling of k , matching the expected $\mathcal{O}(k^2)$ scaling. In contrast, LAMP_{QA} grows by approximately $2\times$, matching the $\mathcal{O}(k)$ checking-layer scaling of the sampled-opening backend. The crossover in absolute constraint count occurs at $k = 256$. At $k = 1,024$, LAMP_{QA} uses 418,149 constraints versus 3,156,298 for Freivalds, a $7.55\times$ reduction. At $k = 4,096$, LAMP_{QA} uses 1,666,315 constraints versus 50,495,818 for Freivalds, a $30.30\times$ reduction.

The same benchmark shows a proving-time crossover at $k = 256$. At $k = 4,096$, LAMP_{QA} proves in 48.54 s versus 405.04 s for Freivalds, an $8.34\times$ speedup. LAMP_{QA} also completes $k = 8,192$ with 3,330,789 constraints and a 151.61 s proving time; Freivalds was not measured at that size in the current artifact. Table 5 breaks down the LAMP_{QA} rows into commitment, SNARK, Merkle, and CP-Link components.

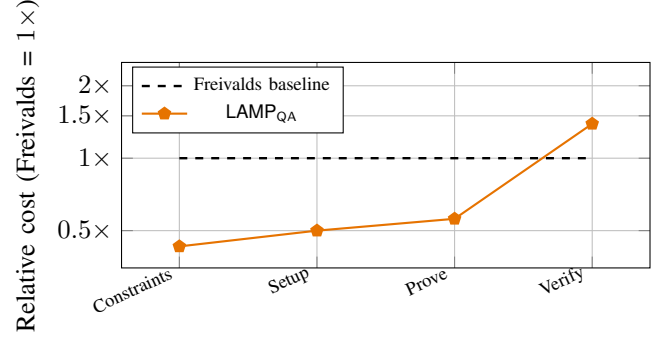


Figure 3. GPT-2 checking-layer case study at sequence length 1,024 with 36 matrix-multiplication claims over input roots. Costs are normalized to the Freivalds baseline.

Verification, online bytes, and GPT-2 workload. LAMP_{QA} verification is nearly constant at about 0.01 s in the matrix benchmark, dominated by the auxiliary linking backend; Freivalds verification is faster because it checks only a direct Groth16 proof. The artifact therefore trades slower verification for smaller prover circuits and faster proving at medium and large dimensions. The LAMP_{QA} online proof grows from 22,404 bytes at $k = 128$ to 59,332 bytes at $k = 4,096$ and 67,268 bytes at $k = 8,192$. At $k = 4,096, t = 128$, publishing the three sampled input-matrix columns alone would require about $3tk \cdot 32 \approx 50$ MB; Appendix C gives detailed byte, setup, and constraint accounting.

Figure 3 and Table 6 report the GPT-2-style workload over input roots. LAMP_{QA} reduces constraints by $2.30\times$ and proving time by $1.78\times$, with lower setup time in this experiment. Verification is $1.39\times$ slower, so latency-sensitive deployments still need batching, amortization, or a more succinct opening/linking backend. The 3.36 MB GPT-2 online proof is larger than the single-matrix rows because it aggregates 36 multiplication claims and their sampled SNARK-side commitments, Merkle authentication paths, and linking material.

Strict parameters and deployment scope. The measured rows use $t = 128$ to study scaling. If $\delta = 1/2$, a conservative per-term 2^{-131} proximity budget uses $t = 133$, increasing sample-dependent work by only $1.0391\times$. At $k = 4,096$, this extrapolates to at most 1.74 million constraints, 50.44 s proving time, and 62 KB online proof bytes. The current artifact does not include measured $t = 133$ rows; these strict numbers are extrapolations from the measured $t = 128$ rows. Online-byte figures exclude raw-matrix encoding, input-root construction, proving/verifying keys, and serialized CRS or generator material. End-to-end deployments preserve the measured prover-time win only if amortized root generation costs stay below the observed margins: 356.50 s for $k = 4,096$ and 178.55 s for the GPT-2 row. The CSV artifact reports point estimates and aggregate setup/proving/verification times; repeated-run variance, confidence intervals, peak memory, and serialized key/CRS/generator sizes remain unmeasured.

$\log k$	Constraints		Prove Time (s)		Verify Time		LAMP _{QA} Online Proof
	LAMP _{QA}	Freivalds	LAMP _{QA}	Freivalds	LAMP _{QA}	Freivalds	
7	54,069	49,610	0.76	0.58	0.01 s	1 ms	22,404 B
8	105,915	197,194	1.52	2.03	0.01 s	1 ms	29,188 B
9	209,995	788,810	3.14	6.89	0.01 s	1 ms	35,204 B
10	418,149	3,156,298	7.19	25.63	0.01 s	2 ms	44,292 B
11	834,075	12,622,154	17.33	102.07	0.01 s	2 ms	51,012 B
12	1,666,315	50,495,818	48.54	405.04	0.01 s	2 ms	59,332 B
13	3,330,789	—	151.61	—	0.01 s	—	67,268 B

TABLE 4. PRIMARY MATRIX-MULTIPLICATION BENCHMARK FOR LAMP_{QA} AND THE OPTIMIZED FREIVALDS SNARK BASELINE. LAMP_{QA} IS THE CURRENT IMPLEMENTATION OF THE SAMPLED-OPENING BACKEND WITH MULTI-COLUMN MERKLE OPENINGS. PROVING TIME EXCLUDES NATIVE MATRIX MULTIPLICATION AND SETUP; THESE ARE RECORDED SEPARATELY IN THE ARTIFACT. THE TABLE MEASURES THE ONLINE MULTIPLICATION-CHECKING LAYER, INCLUDING SAMPLED RS EVALUATIONS, MERKLE OPENINGS, SNARK-SIDE WITNESS COMMITMENTS, AND AUXILIARY OPENING/LINKING CHECKS; ONLINE PROOF BYTES EXCLUDE RAW SAMPLED MATRIX-COLUMN WITNESS VALUES, RAW-MATRIX ENCODING AND INPUT-ROOT CONSTRUCTION FOR CM_A, CM_B, CM_C, AND SETUP/KEY/CRS MATERIAL.

$\log k$	Commit Time	Proving Time			Online Proof Size			Verification Time		
		SNARK	Merkle	CP-Link	SNARK	Merkle	CP-Link	SNARK	Merkle	CP-Link
7	0.10 s	0.63 s	0 ms	0.02 s	260 B	22,016 B	128 B	2 ms	1 ms	3 ms
8	0.29 s	1.18 s	0 ms	0.05 s	260 B	28,800 B	128 B	2 ms	1 ms	3 ms
9	0.91 s	2.17 s	0 ms	0.07 s	260 B	34,816 B	128 B	2 ms	1 ms	3 ms
10	3.20 s	3.88 s	0 ms	0.12 s	260 B	43,904 B	128 B	2 ms	1 ms	3 ms
11	10.02 s	7.10 s	0 ms	0.21 s	260 B	50,624 B	128 B	2 ms	1 ms	3 ms
12	34.57 s	13.55 s	0 ms	0.42 s	260 B	58,944 B	128 B	2 ms	2 ms	3 ms
13	123.87 s	26.81 s	0 ms	0.93 s	260 B	66,880 B	128 B	2 ms	2 ms	3 ms

TABLE 5. LAMP_{QA} COMPONENT BREAKDOWN FROM MEOW_BENCHMARK_RESULTS_V2.CSV. COMMIT TIME IS MATRIXCOMMITTIME PLUS VECTORCOMMITTIME. SNARK PROVING TIME IS CIRCUITPROVETIME; SNARK PROOF SIZE IS THE GROTH16 PROOF SIZE. MERKLE AND CP-LINK COLUMNS REPORT THE CORRESPONDING PROVING, PROOF-SIZE, AND VERIFICATION COMPONENTS FROM THE ARTIFACT.

Metric	LAMP _{QA}	Freivalds	Ratio
Constraints	33.4M	76.8M	0.43×
Setup time	2,078.38 s	4,173.32 s	0.50×
Prove time	230.14 s	408.69 s	0.56×
Verify time	0.68 s	0.49 s	1.39×
LAMP _{QA} online proof	3,355,428 B	—	—

TABLE 6. EXACT GPT-2 CHECKING-LAYER METRICS FROM THE CSV ARTIFACT FOR SEQUENCE LENGTH 1,024, 36 MATRIX-MULTIPLICATION CLAIMS, AND 5 COMMITMENT GROUPS.

7. Conclusion and Related Work

Our work complements verifiable ML systems such as SafetyNets [4], Mystique [2], zkCNN [11], vCNN [12], and zkML [3]; sumcheck/GKR-based VC [13]; dedicated matrix-verification protocols [14]–[17]; and code-based SNARK/proximity systems [5], [8], [18], [19]. Unlike full-circuit ML systems or polynomial-commitment matrix arguments [20], LAMP is a dedicated multiplication checking layer over encoded input roots and is orthogonal to succinct backends, including universal/updatable-SRS and inner-product-style systems [21]–[23]. The security theorem accounts for SNARK soundness, vector/leaf binding, sampled-linking consistency, Freivalds failure, and proximity-sampling error.

The artifact confirms the expected separation. At $k = 4,096$, LAMP_{QA} gives a $30.29\times$ constraint reduction and a $7.24\times$ proving-time speedup over Freivalds-in-SNARK; on the GPT-2 case study it gives a $2.30\times$ constraint reduction and a $1.78\times$ proving-time speedup over input roots. The

remaining limitation is explicit: the artifact measures the online checking layer after input-root generation. Reporting raw-matrix encoding/root costs, key/CRS sizes, reducing verification overhead, and batching many multiplications are the main deployment steps.

Ethics considerations

This work uses local benchmark data and involves no human subjects, live systems, or vulnerability disclosure.

LLM usage considerations

LLMs were used only for editorial assistance; the authors checked all technical claims, proofs, experiments, and citations.

References

- [1] R. Freivalds, “Probabilistic machines can use less running time,” in *Information Processing 77 (Proceedings of IFIP Congress 77)*. Amsterdam, The Netherlands: North-Holland, 1977, pp. 839–842.
- [2] C. Weng, K. Yang, X. Xie, J. Katz, and X. Wang, “Mystique: Efficient conversions for {Zero-Knowledge} proofs with applications to machine learning,” in *30th USENIX Security Symposium (USENIX Security 21)*. Berkeley, CA, USA: USENIX Association, 2021, pp. 501–518.
- [3] B.-J. Chen, S. Waiwitlikhit, I. Stoica, and D. Kang, “Zkml: An optimizing system for ml inference in zero-knowledge proofs,” in *Proceedings of the Nineteenth European Conference on Computer Systems*. New York, NY, USA: ACM, 2024, pp. 560–574.

- [4] Z. Ghodsi, T. Gu, and S. Garg, “Safetynets: Verifiable execution of deep neural networks on an untrusted cloud,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [5] A. Golovnev, J. Lee, S. Setty, J. Thaler, and R. S. Wahby, “Brakedown: Linear-time and field-agnostic snarks for r1cs,” in *Advances in Cryptology – CRYPTO 2023*. Cham: Springer, 2023, pp. 193–226.
- [6] J. T. Schwartz, “Fast probabilistic algorithms for verification of polynomial identities,” *Journal of the ACM*, vol. 27, no. 4, pp. 701–717, 1980.
- [7] R. Zippel, “Probabilistic algorithms for sparse polynomials,” in *Symbolic and Algebraic Computation (EUROSAM ’79)*. Berlin, Heidelberg: Springer, 1979, pp. 216–226.
- [8] T. Xie, Y. Zhang, and D. Song, “Orion: Zero knowledge proof with linear prover time,” in *Advances in Cryptology – CRYPTO 2022*. Cham: Springer, 2022, pp. 299–328.
- [9] C. S. Jutla and A. Roy, “Shorter quasi-adaptive NIZK proofs for linear subspaces,” Cryptology ePrint Archive, Paper 2013/109, 2013. [Online]. Available: <https://eprint.iacr.org/2013/109>
- [10] A. Fiat and A. Shamir, “How to prove yourself: Practical solutions to identification and signature problems,” in *Advances in Cryptology – CRYPTO ’86*. Berlin, Heidelberg: Springer, 1987, pp. 186–194.
- [11] T. Liu, X. Xie, and Y. Zhang, “Zkcnn: Zero knowledge proofs for convolutional neural network predictions and accuracy,” in *Proceedings of the 2021 ACM SIGSAC Conference on Computer and Communications Security*. New York, NY, USA: ACM, 2021, pp. 2968–2985.
- [12] S. Lee, H. Ko, J. Kim, and H. Oh, “vcnn: Verifiable convolutional neural network based on zk-snarks,” *IEEE Transactions on Dependable and Secure Computing*, vol. 21, no. 4, pp. 4254–4270, 2024.
- [13] S. Goldwasser, Y. T. Kalai, and G. N. Rothblum, “Delegating computation: Interactive proofs for muggles,” *Journal of the ACM*, vol. 62, no. 4, pp. 1–64, 2015.
- [14] M. Cong, T. H. Yuen, and S.-M. Yiu, “zkmatrix: Batched short proof for committed matrix multiplication,” in *Proceedings of the 19th ACM Asia Conference on Computer and Communications Security*. New York, NY, USA: ACM, 2024, pp. 289–305.
- [15] M. Cong, T. H. Yuen, and S. M. Yiu, “Dualmatrix: conquering zk-snark for large matrix multiplication,” *Cybersecurity*, vol. 9, no. 1, p. 126, 2026.
- [16] Y. Zhang, M. Zheng, X. Chen, J. Hu, W. Shi, L. Ju, Y. Solihin, and Q. Lou, “zkvc: Fast zero-knowledge proof for private and verifiable computing,” in *2025 62nd ACM/IEEE Design Automation Conference (DAC)*, IEEE. New York, NY, USA: IEEE, 2025, pp. 1–7.
- [17] H. Bennett, K. Gajulapalli, A. Golovnev, and E. Warton, “Matrix multiplication verification using coding theory,” *arXiv preprint arXiv:2309.16176*, 2023.
- [18] S. Ames, C. Hazay, Y. Ishai, and M. Venkitasubramaniam, “Ligero: Lightweight sublinear arguments without a trusted setup,” in *Proceedings of the 2017 ACM SIGSAC Conference on Computer and Communications Security*. New York, NY, USA: ACM, 2017, pp. 2087–2104.
- [19] E. Ben-Sasson, I. Bentov, Y. Horesh, and M. Riabzev, “Fast reed-solomon interactive oracle proofs of proximity,” in *45th International Colloquium on Automata, Languages, and Programming (ICALP 2018)*. Dagstuhl, Germany: Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2018, pp. 14:1–14:17.
- [20] A. Kate, G. M. Zaverucha, and I. Goldberg, “Constant-size commitments to polynomials and their applications,” in *Advances in Cryptology – ASIACRYPT 2010*. Berlin, Heidelberg: Springer, 2010, pp. 177–194.
- [21] M. Maller, S. Bowe, M. Kohlweiss, and S. Meiklejohn, “Sonic: Zero-knowledge snarks from linear-size universal and updatable structured reference strings,” in *Proceedings of the 2019 ACM SIGSAC Conference on Computer and Communications Security*. New York, NY, USA: ACM, 2019, pp. 2111–2128.
- [22] A. Chiesa, Y. Hu, M. Maller, P. Mishra, N. Vesely, and N. Ward, “Marlin: Preprocessing zk-snarks with universal and updatable srs,” in *Advances in Cryptology – EUROCRYPT 2020*. Cham: Springer, 2020, pp. 738–768.
- [23] B. Bünz, J. Bootle, D. Boneh, A. Poelstra, P. Wuille, and G. Maxwell, “Bulletproofs: Short proofs for confidential transactions and more,” in *2018 IEEE Symposium on Security and Privacy*. New York, NY, USA: IEEE, 2018, pp. 315–334.
- [24] T. P. Pedersen, “Non-interactive and information-theoretic secure verifiable secret sharing,” in *Advances in Cryptology – CRYPTO ’91*. Berlin, Heidelberg: Springer, 1991, pp. 129–140.

Appendix A.

Detailed Proof of Theorem 5.1

Proof. We establish completeness and soundness for the interactive protocol, and then explain the Fiat–Shamir variant.

Perfect completeness. Suppose the prover is honest and $\mathbf{AB} = \mathbf{C}$. Let $\mathbf{D} := \mathbf{AB} - \mathbf{C} = \mathbf{0}$. For any challenge vector $\mathbf{r}$, we have $\mathbf{rD} = \mathbf{0}$, so

$$\mathbf{y} = \mathbf{xB} = (\mathbf{rA})\mathbf{B} = \mathbf{r}(\mathbf{AB}) = \mathbf{rC} = \mathbf{z}.$$

For each intermediate vector $v \in \{\mathbf{x}, \mathbf{y}, \mathbf{z}\}$, the honest prover commits to the intermediate vector before I is sampled, commits to the encoded leaf handles, publishes SNARK-side commitments through the CC-SNARK backend, and supplies accepted Merkle openings and CP-Link proofs for both the full vectors and the sampled coordinates. For any queried index $j \in [n]$, linearity of $\mathcal{C}$ gives

$$\text{Fold}(\text{Enc}(\mathbf{A})^{(j)}, \mathbf{r}) = \text{Enc}(\mathbf{rA})_j = \text{Enc}(\mathbf{x})_j,$$

and analogously for the $\mathbf{B}$ and $\mathbf{C}$ folds. Since $\mathbf{y} = \mathbf{z}$, the equality check between their encoded openings also holds at every coordinate. All authenticated openings and backend checks are honest, so the verifier accepts with probability one.

Bad-event decomposition. Let Acc be the event that the verifier accepts a false committed matrix statement, and assume $\mathbf{AB} \neq \mathbf{C}$. Define $\text{Bad}_{\text{SNARK}}$ as acceptance of a false in-circuit SNARK instance, Bad_{cc} as a vector-level or witness-to-leaf linking failure, and Bad_{bind} as two inconsistent accepted leaf-handle openings for the same root and position. For Bad_{cc} , condition on the CC-SNARK binding property: for every accepted SNARK-side vector commitment Γ_v^{vec} it obtains the vector witness v and opening randomness bound to that commitment; for every accepted sampled commitment $\Gamma_{\chi,j}$ it obtains the sampled witness value $a_{\chi,j}$ and opening randomness $\rho_{\chi,j}$ bound to that commitment. The CP-Link batch tag is $T_I = \text{Hash}(\text{FMM-link-batch}, \text{Cm}_{ABC}, \text{Cm}_{XYZ}, I, P_I)$, where

$$P_I = \{(v, \Gamma_v^{\text{vec}}, \mu_v)\}_{v \in \mathcal{W}} \cup \{(\chi, j, \Gamma_{\chi,j}, L_{\chi,j})\}_{\chi,j}.$$

Thus the auxiliary proof is tied to the same vector-level and sampled commitment pairs that appear in the SNARK public statement and the Merkle openings. A vector-level linking failure occurs if the CP-Link proof accepts but, for some $v \in \mathcal{W}$, Γ_v^{vec} and μ_v do not open to the same vector. A sampled linking failure occurs if $\text{MT.Open}(\pi_{\chi,j}, L_{\chi,j}, \text{cm}_{\chi,j}) = 1$

and $\text{CPLink.Verify}(\text{pp}, T_I, P_I, \tau_I) = 1$, but for some sampled pair there is no leaf-opening randomness $\eta_{\chi,j}$ such that $L_{\chi,j}$ opens to the same $a_{\chi,j}$ extracted from the Groth16 witness. For the Pedersen handle instantiation this means there is no $\eta_{\chi,j}$ with

$$L_{\chi,j} = \eta_{\chi,j}H + \sum_{\ell=1}^{d_\chi} (a_{\chi,j})_\ell G_\ell, \\ \Gamma_{\chi,j} - L_{\chi,j} = (\rho_{\chi,j} - \eta_{\chi,j})H.$$

Merkle position binding is charged separately to Bad_{bind} ; outside that event, the accepted $L_{\chi,j}$ is the unique handle at position j under the pre-query root cm_χ . By the assumed backend properties,

$$\Pr[\text{Bad}_{\text{SNARK}}] \leq \epsilon_{\text{SNARK}}(\lambda), \\ \Pr[\text{Bad}_{\text{cc}}] \leq \epsilon_{\text{cc}}(\lambda), \\ \Pr[\text{Bad}_{\text{bind}}] \leq \epsilon_{\text{bind}}(\lambda).$$

Conditioned on the complement of these events, every accepting transcript has a valid SNARK witness, every accepted intermediate vector commitment binds a single vector v before I is sampled, every sampled value used by the Groth16 fold relation is bound to a SNARK-side commitment, every accepted auxiliary opening is linked to that commitment, and every accepted leaf handle is uniquely determined by its root and queried coordinate. The input roots are fixed before the Freivalds challenge is sampled, and the intermediate commitments are fixed before the sampled query tuple is sampled.

Freivalds randomization. Under the same conditioning, the matrices bound to $\text{cm}_A, \text{cm}_B, \text{cm}_C$ are fixed and satisfy $\mathbf{AB} \neq \mathbf{C}$. Let $\mathbf{D} = \mathbf{AB} - \mathbf{C} \neq 0$. There exists a column $\mathbf{d}_j$ of $\mathbf{D}$ with $\mathbf{d}_j \neq 0$. The condition $\mathbf{rD} = 0$ implies

$$P_j(r) := \langle \mathbf{r}, \mathbf{d}_j \rangle = \sum_{i=0}^{k-1} (\mathbf{D})_{i,j} r^i = 0.$$

Since P_j is a nonzero polynomial of degree at most $k-1$ and r is sampled after the input commitments are fixed, Lemma 2.2 gives

$$\Pr_{r \leftarrow \mathbb{F}}[\mathbf{rD} = 0] \leq \Pr_{r \leftarrow \mathbb{F}}[P_j(r) = 0] \leq \frac{k-1}{|\mathbb{F}|}.$$

Proximity sampling. Now condition further on $\mathbf{rD} \neq 0$. No vectors can simultaneously satisfy all four global relations:

$$\mathbf{x} = \mathbf{rA}, \quad \mathbf{y} = \mathbf{xB}, \quad \mathbf{z} = \mathbf{rC}, \quad \mathbf{y} = \mathbf{z}.$$

Otherwise the four equalities would imply $\mathbf{rAB} = \mathbf{rC}$, contradicting $\mathbf{rD} \neq 0$. Thus at least one relation is false. Consider $\mathbf{x} \neq \mathbf{rA}$; the other cases are identical. Injectivity and distance of the code imply that the two codewords $\text{Enc}(\mathbf{x})$ and $\text{Enc}(\mathbf{rA})$ disagree on at least δn coordinates. By code linearity, the sampled fold value from the committed rows of $\mathbf{A}$ equals $\text{Enc}(\mathbf{rA})_j$ at coordinate j . Since the

sampled column used by Groth16 is bound to the authenticated input leaf handle, and the sampled vector value is both bound to its authenticated leaf handle and checked to equal the circuit-computed codeword $\text{Enc}(\mathbf{x})_j$, the fold check rejects whenever j is in the disagreement set. A single uniformly sampled coordinate misses this disagreement set with probability at most $1 - \delta$, and t independent sampled coordinates all miss it with probability at most $(1 - \delta)^t$. Union-bounding over the four possible violated relations yields $4(1 - \delta)^t$.

Combining the bounds and parameters. Summing the backend, Freivalds, and proximity errors gives

$$\Pr[\text{Acc} \mid \mathbf{AB} \neq \mathbf{C}] \leq \epsilon_{\text{SNARK}}(\lambda) + \epsilon_{\text{cc}}(\lambda) + \epsilon_{\text{bind}}(\lambda) \\ + \frac{k-1}{|\mathbb{F}|} + 4(1 - \delta)^t.$$

For $\lambda = 128$, a 256-bit field makes $(k-1)/|\mathbb{F}|$ negligible for the reported dimensions. The proximity term is at most $2^{-\lambda}$ when $t = \lceil (\lambda+2)/\log_2(1/(1-\delta)) \rceil$; for example, for a code with $\delta = 1/2$, $t = 130$ suffices for the proximity term alone. For a conservative aggregate 128-bit budget over all terms, one should allocate the budget across terms, e.g., by targeting each term at 2^{-131} .

Fiat-Shamir variant. The non-interactive transcript derives both r and I from domain-separated hashes of the transcript prefix, including the public input roots, intermediate vector commitments, and encoded-vector roots. In the random oracle model, the preceding public-coin proof applies with the standard Fiat-Shamir loss $\epsilon_{\text{FS}}(q_H)$ for an adversary that makes q_H random-oracle queries. This proves Theorem 5.1. $\square$

Appendix B. CP-Link Backend Instantiation

The artifact uses Pedersen-style [24] commitments

$$\Gamma = \rho_\Gamma H + \sum_{\ell=1}^{d_\chi} a_\ell G_\ell, \quad C = \rho_C H + \sum_{\ell=1}^{d_\chi} a_\ell G_\ell,$$

with $d_\chi = k$ for matrix columns and full intermediate vectors, and $d_\chi = 1$ for vector coordinates. For vector-level links, C is the pre-query vector commitment μ_v . For sampled links, C is the authenticated Merkle leaf handle $L_{\chi,j}$. Leaf handles are authenticated by the Merkle tree, while the CC-SNARK backend binds Γ to the vector or sampled witness value used by the circuit. The link proof shows that the SNARK-side commitment and the external commitment open to the same object.

Theorem B.1 (Batched CP-Link backend). *Assume the Pedersen handle parameters are generated so that the handle commitment is binding, the QA-NIZK link proof for the linear batch relation is sound, and the batching challenge is sampled by the Fiat-Shamir random oracle after the batch statement is fixed. For $Q = |S|t + |W|$ vector-level and sampled objects, the batched CP-Link verifier accepts a false batch with probability at most $\epsilon_{\text{ped}} + \epsilon_{\text{QA}} + Q/|\mathbb{F}|$.*

The artifact serializes the Q logical relations as one online batched proof; generator and CRS bytes are setup material and are not included in the online proof-size claim.

The logical CP-Link relation contains one equation per vector-level or sampled object:

$$\Gamma_q - C_q = \Delta_q H, \quad q \in [Q],$$

where $Q = |\mathcal{S}|t + |\mathcal{W}|$ and q indexes both vector-level pairs $(v, \Gamma_v^{\text{vec}}, \mu_v)$ and sampled pairs $(\chi, j, \Gamma_{\chi,j}, L_{\chi,j})$. The batch statement is

$$P_I = \{(\text{tag}_q, \Gamma_q, C_q)\}_{q=1}^Q, \\ T_I = \text{Hash}(\text{FMM-link-batch}, \text{Cm}_{ABC}, \text{Cm}_{XYZ}, I, P_I).$$

The implementation batches these equations by deriving $\beta \leftarrow \text{Hash}(\text{FMM-link-rand}, T_I, \{\Gamma_q, C_q\}_{q=1}^Q)$ and checking a single randomized linear relation

$$\sum_{q=1}^Q \beta^q (\Gamma_q - C_q) = \left(\sum_{q=1}^Q \beta^q \Delta_q \right) H.$$

If any individual equality is false, the left-minus-right side is a nonzero degree- Q polynomial in β , so the batch can pass with probability at most $Q/|\mathbb{F}|$ in the random-oracle model. Combining this with Pedersen binding and QA-NIZK soundness gives the bound in Theorem B.1. The QA-NIZK statement is $(\text{pp}, T_I, \{\Gamma_q, C_q\}_{q=1}^Q)$ and the witness is the aggregate difference opening $\sum_q \beta^q \Delta_q$. Its CRS fixes the handle generators and linear-relation verification parameters; if serialized, those bytes are setup/key/CRS material rather than online proof bytes. Thus the 128-byte value reported in the experiments is one online batched QA-NIZK proof for all vector-level and sampled links, not a per-link cost.

Appendix C. Additional Evaluation Details

This appendix records the measurement boundary behind the main evaluation. Tables 7 and 8 first state the fixed experimental parameters and which costs are covered by the artifact. Tables 9–12 then separate serialized proof material from private Groth16 witness data and from costs that remain unmeasured, such as raw-matrix encoding/root construction and serialized key/CRS size. The intent is to make the main text’s proof-size and speedup claims auditable without treating the checking-layer measurements as full end-to-end VC costs.

Table 7 fixes the common parameter choices used by the reported rows. In particular, the measured rows use $t = 128$ for scaling, while the table also records the nearby query counts needed for stricter proximity-only and per-term security budgets.

Parameter	Value in main experiments
Field size	256-bit prime field
Server	AMD EPYC 7B13, 32 cores, 251.69 GiB RAM, Linux/amd64
Matrix dimensions	LAMP _{QA} : 2^7 – 2^{13} , Freivalds: 2^7 – 2^{12}
Code rate	$\rho = 1/2$
Codeword length	$n = 2k$
Sampled positions	$t = 128$
Primary linking layer	Pre-query vector commitments, SNARK-side commitments, QA-NIZK/CP-Link, and multi-column Merkle openings
Backend model	CP-Link sampled intermediate backend; input roots are fixed before r
Query-count targets	Measured rows use $t = 128$; proximity-only 2^{-128} at $\delta = 1/2$ uses $t = 130$; conservative per-term 2^{-131} uses $t = 133$.

TABLE 7. CONCRETE PARAMETERS USED IN THE MAIN EXPERIMENTS.

Table 8 is a measurement-scope checklist rather than a performance result. It distinguishes quantities measured directly by the CSV artifact from costs that remain outside the online proof bytes, most notably raw-matrix encoding/root construction and serialized key/CRS sizes.

Metric	Covered?	Interpretation
Constraint count	Yes, aggregate	Groth16 circuit size including sampled RS evaluations and fold checks; per-gadget constraint counters are not reported, and CC-SNARK/linking overhead is measured in time/proof bytes rather than as extra R1CS constraints.
Online proving time	Yes	Excludes native matrix multiplication and one-time setup.
Verification time	Yes	Includes current QA-NIZK/CP-Link and opening verification overhead.
Proof size	Yes	Includes Groth16, compact sampled commitments/openings, and QA-NIZK material; raw sampled matrix columns are Groth16 witness data, not serialized proof bytes.
Setup time	Partial	Recorded for the matrix and GPT-2 workloads; setup/key sizes are not reported.
Input-root generation	Partial	Root construction is outside online proof bytes; raw-matrix encoding/root costs must be added for end-to-end deployments.
Peak memory and variance	No	Needed for a full systems artifact evaluation.

TABLE 8. MEASUREMENT COVERAGE IN THE CURRENT ARTIFACT.

Table 9 explains the proof-size convention used in the main text. The sampled matrix columns are large, but they are private Groth16 witness data; the serialized proof contains only compact handles, authentication paths, the Groth16 proof, and the batched CP-Link material.

Object	Serialized?	Where it is checked / charged
Input column $a_{M,j}$	No	k field elements are Groth16 witness data; handles and paths are serialized.
Vector leaf $a_{v,j}$	No	One field element is used as Groth16 witness data and checked by $a_{v,j} = \text{Enc}(v)_j$; the sampled SNARK-side commitment is linked to the Merkle leaf handle.
Merkle path	Yes	Authenticates $L_{\chi,j}$ in AuxV; charged to proof bytes and verify time.
CP-Link proof	Yes	One batched QA-NIZK proof τ_I links all vector-level and sampled commitment pairs.
Groth16 proof	Yes	Verifies sampled RS evaluations and fold/equality checks, while its CC-SNARK interface binds full vector witnesses and sampled values to SNARK-side commitments.

TABLE 9. PROOF-MATERIAL ACCOUNTING FOR THE CP-LINK SAMPLED-OPENING BACKEND.

Table 10 instantiates the accounting for two representative matrix sizes. The reported totals are the serialized online proof bytes in the artifact; roots and vector commitments are public inputs, and raw sampled columns are intentionally not serialized.

Case	Opening payload	Online proof	Note
$k = 1,024$	43,904 B	44,292 B	Adds 260 B Groth16 and 128 B batched CP-Link; roots/vector commitments are public inputs.
$k = 4,096$	58,944 B	59,332 B	Raw sampled columns would be 50.3 MB if serialized.

TABLE 10. REPRESENTATIVE BYTE-LEVEL ACCOUNTING FROM MEOW_BENCHMARK_RESULTS_V2.CSV.

Table 11 records setup time for representative rows and separates protocol, circuit, and CP-Link setup components. The current artifact does not report serialized proving/verifying-key or CRS bytes; this is why the main text states proof size as online proof material excluding key/CRS data.

Case	Total setup	Protocol	Circuit	CP-Link	Key / CRS size
$k = 1,024$	33.16 s	0.18 s	31.83 s	1.15 s	Not reported by the current artifact.
$k = 4,096$	130.08 s	0.71 s	125.01 s	4.37 s	Not reported by the current artifact.

TABLE 11. SETUP ACCOUNTING FOR REPRESENTATIVE ROWS.

Table 12 separates the dominant deterministic $3tk$ fold terms from the remaining checked constraints. The residual includes sampled RS evaluations, equality checks, and circuit glue, so the table is a structural audit of the aggregate constraint counts rather than a per-gadget profiler.

Case	Total	$3tk$ folds	Residual	Raw cols.
$k = 1,024$	418,149	393,216	24,933	12.6 MB
$k = 4,096$	1,666,315	1,572,864	93,451	50.3 MB

TABLE 12. STRUCTURAL CONSTRAINT AND RAW-COLUMN-SIZE ACCOUNTING FOR TWO REPRESENTATIVE ROWS.

The accounting tables should be read as lower-level support for the main evaluation, not as additional benchmark claims. They justify three statements used throughout the paper: the large sampled columns remain private witness data, the serialized online proof is dominated by compact opening material rather than raw columns, and the current artifact still leaves key/CRS size and raw-matrix encoding/root construction as deployment costs outside the online proof-byte measurements.